

Configure scala on Jupyter.

Option I: Using spylon – Execute Scala spark program.

```
#pip3 install spylon-kernel  
#python3 -m spylon_kernel install  
#export SPARK_HOME=/opt/spark  
#jupyter notebook
```

In the web console, select New -> Spylon-kernel and enter the following commands.

The screenshot shows a Jupyter Notebook interface. At the top, there's a toolbar with File, Edit, View, Insert, Cell, Kernel, Help, and a Trusted button. To the right of the toolbar is a user icon and a Logout button. A red circle highlights the 'spylon-kernel' option in the Kernel dropdown menu. The main workspace shows two code cells:

```
In [1]: 1+1
Initialization Scala interpreter ...
Spark Web UI available at http://spark0:4040
SparkContext available as 'sc' (version = 3.1.2, master = local[*], app id = local-1645081006872)
SparkSession available as 'spark'

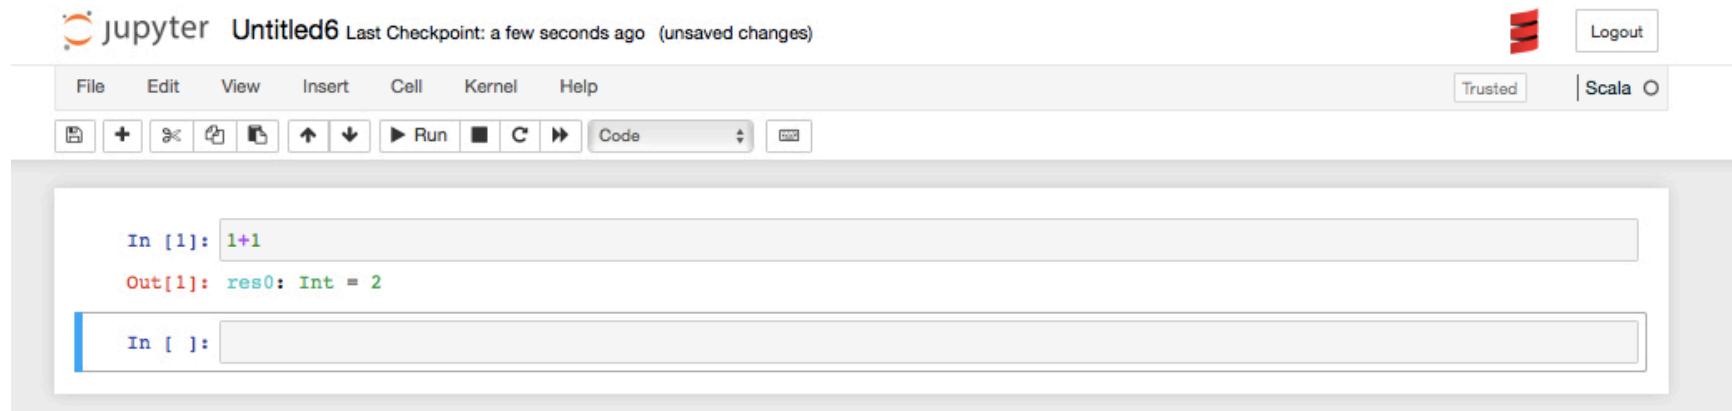
Out[1]: res0: Int = 2

In [2]: sc
Out[2]: res1: org.apache.spark.SparkContext = org.apache.spark.SparkContext@17d3a53d
```

The first cell runs `1+1` and outputs `res0: Int = 2`. The second cell runs `sc` and outputs `res1: org.apache.spark.SparkContext = org.apache.spark.SparkContext@17d3a53d`.

Option 2: using Scala in Jupyter

<https://almond.sh/docs/quick-start-install>



Or using Anaconda: (Optional)

Anaconda 4.2.0

For Linux

Anaconda is BSD licensed which gives you permission to use Anaconda commercially and for redistribution.

[Changelog](#)

Download the installer

Optional: Verify data integrity with [MD5 or SHA-256](#) [More info](#)

In your terminal window type one of the below and follow the instructions:

Python 3.5 version

bash Anaconda3-4.2.0-Linux-x86\_64.sh

Python 2.7 version

```
bash Anaconda2-4.2.0-Linux-x86_64.sh
```

NOTE: Include the "bash" command even if you are not using the bash shell.

```
[root@tos Software]# bash  
[root@tos Software]# sh Anaconda3-4.2.0-Linux-x86_64.sh  
  
Welcome to Anaconda3 4.2.0 (by Continuum Analytics, Inc.)  
  
In order to continue the installation process, please review the license  
agreement.  
Please, press ENTER to continue  
>>> [REDACTED]
```

```
Anaconda3 will now be installed into this location:  
/root/anaconda3
```

- Press ENTER to confirm the location
- Press CTRL-C to abort the installation
- Or specify a different location below

```
[/root/anaconda3] >>> /anaconda3 [REDACTED]  
  
Do you wish the installer to prepend the Anaconda3 install location  
to PATH in your /root/.bashrc ? [yes|no]  
[no] >>> yes
```

```
Prepending PATH=/anaconda3/bin to PATH in /root/.bashrc  
A backup will be made to: /root/.bashrc-anaconda3.bak
```

```
For this change to become active, you have to open a new terminal.
```

```
Thank you for installing Anaconda3!
```

```
Share your notebooks and packages on Anaconda Cloud!  
Sign up for free: https://anaconda.org
```

```
[root@tos Software]# [REDACTED]
```

```
#bash
```

Run the following command to install Jupyter.

```
#conda install jupyter
```

```
[root@tos Software]# conda install jupyter
Fetching package metadata .....
Solving package specifications: .....

Package plan for installation in environment /anaconda3:

The following packages will be downloaded:

  package          |      build
  -----|-----
  conda-env-2.6.0   |          0      502 B
  requests-2.12.4    |      py35_0     800 KB
  pyopenssl-16.2.0   |      py35_0      70 KB
  conda-4.3.7        |      py35_0     491 KB
  -----
                           Total:     1.3 MB
```

```
The following packages will be UPDATED:

conda:      4.2.9-py35_0 --> 4.3.7-py35_0
pyopenssl: 16.0.0-py35_0 --> 16.2.0-py35_0
requests:   2.11.1-py35_0 --> 2.12.4-py35_0

Proceed ([y]/n)? y

Fetching packages ...
conda-env-2.6. 100% [########################################| Time: 0:00:00 129.08 kB/s
requests-2.12. 100% [########################################| Time: 0:00:11 70.63 kB/s
pyopenssl-16.2 100% [########################################| Time: 0:00:01 53.66 kB/s
conda-4.3.7-py 100% [########################################| Time: 0:00:04 104.02 kB/s
Extracting packages ...
[    COMPLETE     ]|########################################| 100%
Unlinking packages ...
[    COMPLETE     ]|########################################| 100%
Linking packages ...
[    COMPLETE     ]|########################################| 100%
dbus post-link :: /etc/machine-id not found ..
dbus post-link :: .. using /proc/sys/kernel/random/boot_id
[root@tos Software] #
```

**Note:**

Configure default Python version, If you have multiples python version

#alternatives --config python3

or

#update-alternatives --config python3

Reference:

<https://www.sicara.ai/blog/2017-05-02-get-started-pyspark-jupyter-notebook-3-minutes>

<https://community.hortonworks.com/articles/75551/installing-and-exploring-spark-20-with-jupyter-not.html>

<https://medium.com/@bogdan.cojocar/how-to-run-scala-and-spark-in-the-jupyter-notebook-328a80090b3b>

-----Lab Ends Here -----

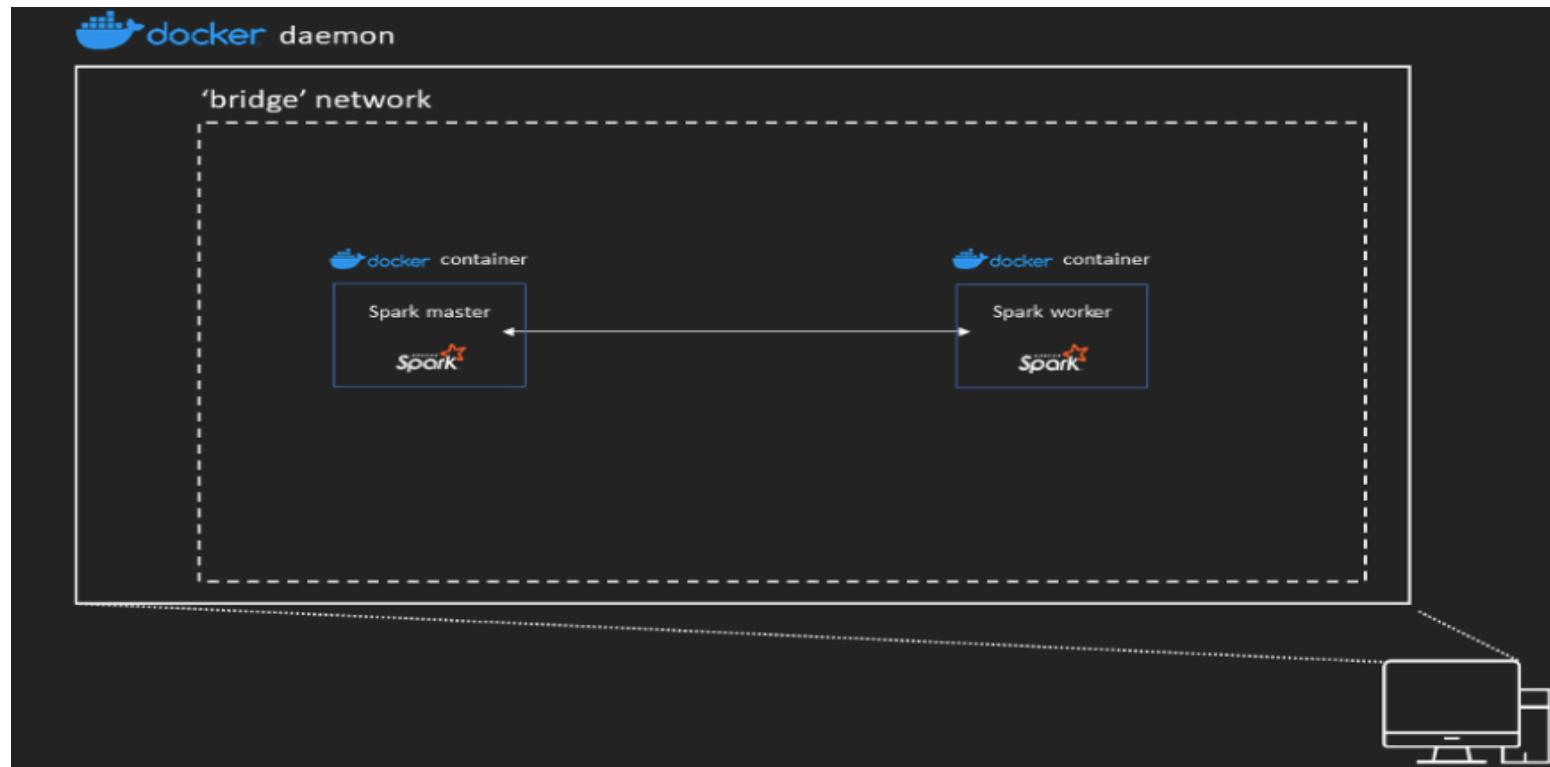
## 15. Explore Spark Two Nodes Cluster – 90 Minutes

In this lab, we will configure 2 nodes spark standalone cluster and deploy application on the cluster.

Prerequisites:

- You required 2 instances of node.
  - o Spark0
  - o Spark1

Architecture



Roles	Hostname	Remarks
<b>master</b>	Sparko	Install Spark
<b>slave/worker</b>	Spark1	Install Spark

Let us start the Spark master – **sparko**.

Install spark on the master node. You can refer the first lab or skip already done.

Execute the following in Sparko

Setup a Spark master node

Change the Master Port to 8090, there are some issues when it is executed in 7077 port in a docker environment.

Execute on the **sparko** – Terminal.

```
#export PATH=$PATH:/opt/spark/bin
#export SPARK_MASTER_PORT=8090
#spark-class org.apache.spark.deploy.master.Master
```

```
@5d486484cd1c:/software (docker)          361           bash          362
20/09/12 02:33:29 INFO Master: Started daemon with process name: 35@5d486484cd1c
20/09/12 02:33:29 INFO SignalUtils: Registered signal handler for TERM
20/09/12 02:33:29 INFO SignalUtils: Registered signal handler for HUP
20/09/12 02:33:29 INFO SignalUtils: Registered signal handler for INT
WARNING: An illegal reflective access operation has occurred
WARNING: Illegal reflective access by org.apache.spark.unsafe.Platform (file:/opt/spark/jars/spark-unsafe_2.12-3.0.1.jar) to constructor java.nio.DirectByteBuffer(long,int)
WARNING: Please consider reporting this to the maintainers of org.apache.spark.unsafe.Platform
WARNING: Use --illegal-access=warn to enable warnings of further illegal reflective access operations
WARNING: All illegal access operations will be denied in a future release
20/09/12 02:33:30 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
20/09/12 02:33:30 INFO SecurityManager: Changing view acls to: root
20/09/12 02:33:30 INFO SecurityManager: Changing modify acls to: root
20/09/12 02:33:30 INFO SecurityManager: Changing view acls groups to:
20/09/12 02:33:30 INFO SecurityManager: Changing modify acls groups to:
20/09/12 02:33:30 INFO SecurityManager: SecurityManager: authentication disabled; ui acls disabled; users with view permissions: Set(root); groups with view permissions: Set(); users with modify permissions: Set(root); groups with modify permissions: Set()
20/09/12 02:33:31 INFO Utils: Successfully started service 'sparkMaster' on port 7077.
20/09/12 02:33:31 INFO Master: Starting Spark master at spark://172.17.0.2:7077
20/09/12 02:33:31 INFO Master: Running Spark version 3.0.1
20/09/12 02:33:32 INFO Utils: Successfully started service 'MasterUI' on port 8080.
20/09/12 02:33:32 INFO MasterWebUI: Bound MasterWebUI to 0.0.0.0, and started at http://5d486484cd1c:8080
20/09/12 02:33:32 INFO Master: I have been elected leader! New state: ALIVE
```

Start worker process On **sparko**: In a separate terminal.

Attach a worker node to the cluster, execute the following in the sparko container.

```
#export PATH=$PATH:/opt/spark/bin  
#spark-class org.apache.spark.deploy.worker.Worker -c 1 -m 1G spark://sparko:8090
```

```
20/09/12 02:55:17 INFO Worker: Starting Spark worker 172.17.0.2:38483 with 1 cores, 2.0 GiB RAM  
20/09/12 02:55:17 INFO Worker: Running Spark version 3.0.1  
20/09/12 02:55:17 INFO Worker: Spark home: /opt/spark  
20/09/12 02:55:17 INFO ResourceUtils: =====  
20/09/12 02:55:17 INFO ResourceUtils: Resources for spark.worker:  
  
20/09/12 02:55:17 INFO ResourceUtils: =====  
20/09/12 02:55:17 INFO Utils: Successfully started service 'WorkerUI' on port 8081.  
20/09/12 02:55:17 INFO WorkerWebUI: Bound WorkerWebUI to 0.0.0.0, and started at http://224e200bfc56:8081  
20/09/12 02:55:17 INFO Worker: Connecting to master 172.17.0.2:7077...  
20/09/12 02:55:18 INFO TransportClientFactory: Successfully created connection to /172.17.0.2:7077 after 40 ms (0 ms spent in bootstraps)  
20/09/12 02:55:18 INFO Worker: Successfully registered with master spark://172.17.0.2:7077  
20/09/12 02:55:18 INFO Worker: Asked to launch executor app-20200912025310-0000/0 for Spark shell  
20/09/12 02:55:18 INFO SecurityManager: Changing view acls to: root  
20/09/12 02:55:18 INFO SecurityManager: Changing modify acls to: root  
20/09/12 02:55:18 INFO SecurityManager: Changing view acls groups to:  
20/09/12 02:55:18 INFO SecurityManager: Changing modify acls groups to:  
20/09/12 02:55:18 INFO SecurityManager: SecurityManager: authentication disabled; ui acls disabled; users with view permissions: Set (root); groups with view permissions: Set(); users with modify permissions: Set(root); groups with modify permissions: Set()  
20/09/12 02:55:18 INFO ExecutorRunner: Launch command: "/opt/java/bin/java" "-cp" "/opt/spark/conf/:/opt/spark/jars/*" "-Xmx1024M" "-Dspark.driver.port=35927" "org.apache.spark.executor.CoarseGrainedExecutorBackend" "--driver-url" "spark://CoarseGrainedScheduler@224e200bfc56:35927" "--executor-id" "0" "--hostname" "172.17.0.2" "--cores" "1" "--app-id" "app-20200912025310-0000" "--worker-url" "spark://Worker@172.17.0.2:38483"
```

If unable to connect to **sparko** replace it with the container IP.

Refresh the web UI, ensure that you can see a worker as shown below.

<http://127.0.0.1:8080>

The screenshot shows the Apache Spark 3.0.1 master web UI. At the top, it displays "Spark Master at spark://172.17.0.2:8090". Below this, various system metrics are listed:

- URL: spark://172.17.0.2:8090
- Alive Workers: 1
- Cores in use: 1 Total, 0 Used
- Memory in use: 1024.0 MIB Total, 0.0 B Used
- Resources in use:
- Applications: 0 Running, 0 Completed
- Drivers: 0 Running, 0 Completed
- Status: ALIVE

Under the "Workers (1)" section, there is a table with the following data:

Worker Id	Address	State	Cores	Memory	Resources
worker-20210118134218-172.17.0.2-34667	172.17.0.2:34667	ALIVE	1 (0 Used)	1024.0 MIB (0.0 B Used)	

Below the workers table, there are sections for "Running Applications (0)" and "Completed Applications (0)", each with its own table header.

IP against the worker node will depends on the environment or your system.

## Start the Node on spark1

Ensure that you have install spark in the **Spark1** too.

On the **spark1** terminal, it will register to the spark master – **sparko**

```
# spark-class org.apache.spark.deploy.worker.Worker -c 1 -m 1G spark://sparko:8090
```

At the end of this step, you should have 2 worker nodes as shown below: Refresh the Spark Web UI.

The screenshot shows the Spark Web UI interface. At the top, it displays "Spark 3.0.1" and "Spark Master at spark://172.19.0.3:8090". Below this, there are several status metrics:

- URL: spark://172.19.0.3:8090
- Alive Workers: 2
- Cores in use: 2 Total, 0 Used
- Memory in use: 2.0 GiB Total, 0.0 B Used
- Resources in use:
- Applications: 0 Running, 0 Completed
- Drivers: 0 Running, 0 Completed
- Status: ALIVE

Below these metrics, there is a section titled "Workers (2)" which lists the two registered workers:

Worker Id	Address	State	Cores	Memory	Resources
worker-20200912142723-172.19.0.3-35677	172.19.0.3:35677	ALIVE	1 (0 Used)	1024.0 MiB (0.0 B Used)	
worker-20200912142924-172.19.0.2-38199	172.19.0.2:38199	ALIVE	1 (0 Used)	1024.0 MiB (0.0 B Used)	

There are also sections for "Running Applications (0)" and "Completed Applications (0)", each with an empty table.

Open a pyspark shell and connect to the Spark cluster

Let us execute Spark shell using Cluster mode:

You can perform from any client/server node i.e windows desktop client.

## Connecting an Application to the Cluster

To run an application on the Spark cluster, simply pass the `spark://IP:PORT` URL of the master as to the `SparkContext` constructor.

To run an interactive Spark shell against the cluster(Specify the master IP), run the following command from the bin folder:

We are executing from the slave node – **spark1** console.

```
#pyspark --conf spark.executor.memory=512mb --conf spark.executor.cores=1 --master spark://sparko:8090
```

```
[root@spark0 /]# pyspark --conf spark.executor.memory=512mb --conf spark.executor.cores=1 --master spark://spark0:8090
Python 3.6.8 (default, Sep 10 2021, 09:21:56)
[GCC 8.5.0 20210514 (Red Hat 8.5.0-3)] on linux
Type "help", "copyright", "credits" or "license" for more information.
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
23/07/03 10:10:49 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Welcome to

    ___
   / _ \_ ___ _ _ / _ \_
  _\ \ V \ - \ \_ \ \_ / ' \_
 /__ / .__\ \_, \_ / / \_\ \_ version 3.3.0
   /_/

Using Python version 3.6.8 (default, Sep 10 2021 09:21:56)
Spark context Web UI available at http://spark0:4040
Spark context available as 'sc' (master = spark://spark0:8090, app id = app-20230703101049-0000).
SparkSession available as 'spark'.
>>> █
```

You can verify the application execution from the web UI. There should be an entry in the Applications section.

**Spark Master at spark://192.168.188.173:7077**

URL: spark://192.168.188.173:7077  
 REST URL: spark://192.168.188.173:6066 (cluster mode)  
 Alive Workers: 2  
 Cores in use: 2 Total, 2 Used  
 Memory in use: 2.0 GB Total, 2.0 GB Used  
 Applications: 1 Running, 0 Completed  
 Drivers: 0 Running, 0 Completed  
 Status: ALIVE

**Workers**

Worker Id	Address	State	Cores	Memory
worker-20170305001707-192.168.188.174-35897	192.168.188.174:35897	ALIVE	1 (1 Used)	1024.0 MB (1024.0 MB Used)
worker-20170305001728-192.168.188.173-40860	192.168.188.173:40860	ALIVE	1 (1 Used)	1024.0 MB (1024.0 MB Used)

**Running Applications**

Application ID	Name	Cores	Memory per Node	Submitted Time	User	State	Duration
app-20170305004646-0000	(kill) Spark shell	2	1024.0 MB	2017/03/05 00:46:46	root	RUNNING	6.4 min

Let's make a new RDD from the text of the README file in the Spark source directory:

```
textFile = sc.textFile("/opt/spark/README.md")
```

**Let's start with a few actions:**

```
textFile.count() #Number of items in this RDD
textFile.first() # First item in this RDD
```

```
>>> textFile = sc.textFile("/opt/spark/README.md")
>>> textFile.count() #Number of items in this RDD
124
>>> textFile.first() # First item in this RDD
'# Apache Spark'
>>> █
```

Now let's use a transformation. We will use the `filter` transformation to return a new RDD with a subset of the items in the file. Use `collect` to display the output.

```
linesWithSpark = textFile.filter(lambda line : "Spark" in line)
linesWithSpark.collect()
```

```
>>> linesWithSpark = textFile.filter(lambda line : "Spark" in line)
>>> linesWithSpark.collect()
['# Apache Spark', 'Spark is a unified analytics engine for large-scale data processing. It provides', 'rich set of higher-level tools including Spark SQL for SQL and DataFrames,', 'pandas API on Spark for pandas workloads, MLlib for machine learning, GraphX for graph processing,', '[![PySpark Coverage](https://codecov.io/gh/apache/spark/branch/master/graph/badge.svg)](https://codecov.io/gh/apache/spark)', 'You can find the latest Spark documentation, including a programming', '## Building Spark', 'Spark is built using [Apache Maven](https://maven.apache.org/).', 'To build Spark and its example programs, run:', '["Building Spark"](/spark.apache.org/docs/latest/building-spark.html).', 'For general development tips, including info on developing Spark using an IDE, see ["Useful Developer Tools"](/spark.apache.org/developer-tools.html).', 'The easiest way to start using Spark is through the Scala shell:', 'Spark also comes with several sample programs in the `examples` directory.', './bin/run-example SparkPi', 'MASTER=spark://host:7077 ./bin/run-example Spark Pi', 'Testing first requires [building Spark](#building-spark). Once Spark is built, tests', 'Spark uses the Hadoop core library to talk to HDFS and other Hadoop-supported', 'Hadoop, you must build Spark against the same version that your cluster runs.', 'in the online documentation for an overview on how to configure Spark.', 'Please review the [Contribution to Spark guide](https://spark.apache.org/contributing.html)']
```

We can chain together transformations and actions:

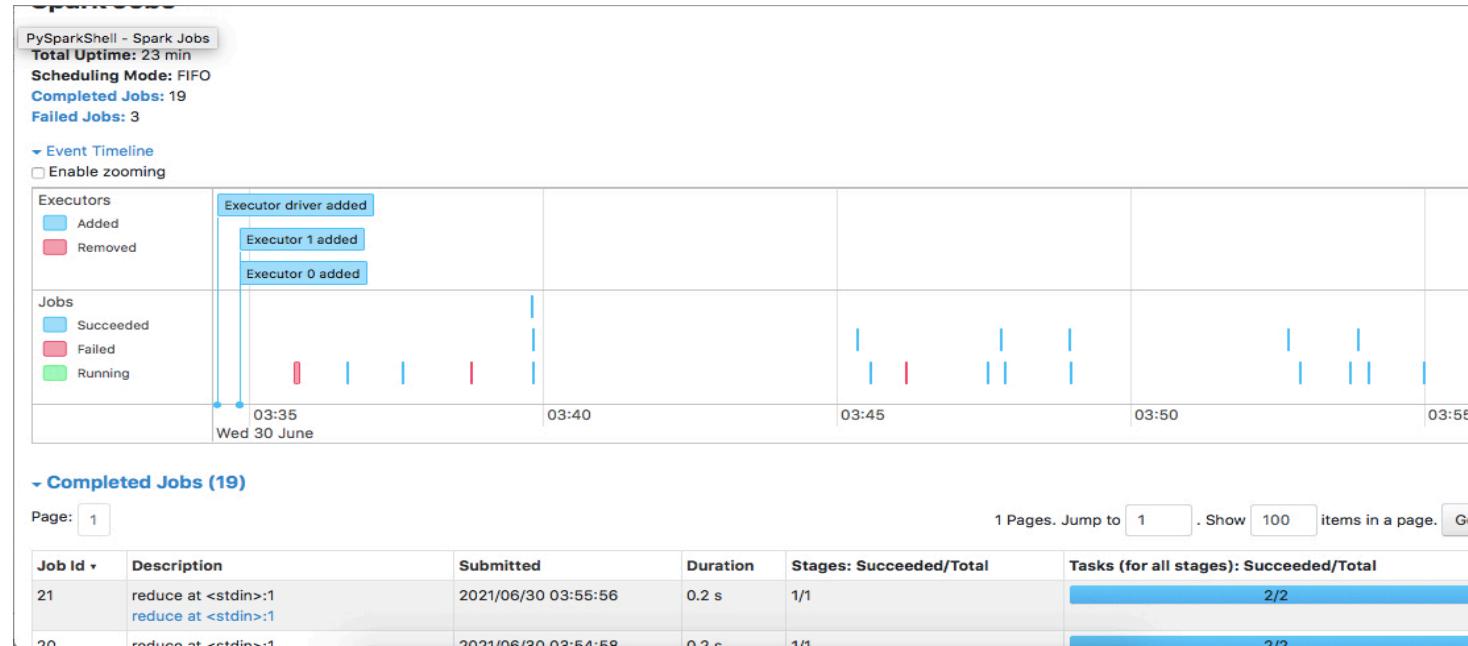
```
textFile.filter(lambda line : "Spark" in line).count() # How many lines contain "Spark"?
```

Let's say we want to find the line with the most words

```
textFile.map(lambda line : len(line.split(" "))).reduce(lambda a, b: a if (a > b) else b)
```

```
>>> textFile.filter(lambda line : "Spark" in line).count() # How many lines contain "Spark"?  
20  
>>> textFile.map(lambda line : len(line.split(" "))).reduce(lambda a, b: a if (a > b) else b)  
16  
>>> 
```

Refer the web UI for the Job details as shown below.



The console should display the result as shown above.

Hints: (For submitting job in cluster)

You need to complete the **Lab - Explore PySpark Application** before going ahead. Submit from the /software folder, since the input file exist in that folder.

You need to copy the data files in both the nodes. However, for our lab its already present in the spark folder. Execute the command from the folder where the python code is.

```
#spark-submit --master spark://spark0:8090 SimpleApp.py --conf spark.executor.memory=256mb
```

Verify the output.

```
21/06/30 04:11:32 INFO TaskSchedulerImpl: Killing all running tasks in stage 1: Stage finished
21/06/30 04:11:32 INFO DAGScheduler: Job 1 finished: count at /opt/data/SimpleApp.py:9, took 0.174200 s
Lines with a: 64, lines with b: 32
21/06/30 04:11:32 INFO SparkUI: Stopped Spark web UI at http://spark0:4041
21/06/30 04:11:32 INFO MapOutputTrackerMasterEndpoint: MapOutputTrackerMasterEndpoint stopped!
21/06/30 04:11:32 INFO MemoryStore: MemoryStore cleared
```

----- Lab Ends Here -----

<https://towardsdatascience.com/diy-apache-spark-docker-bb4f11c10d24>

## 16. Explore Spark on Hadoop YARN – 150 Minutes

In this lab, you will deploy Spark Application on Hadoop Cluster.

Steps:

- Configure and install YARN on hadoopo.
- One Node will be exclusively running YARN cluster
- Existing sparko node to be used to submit spark job on YARN cluster.

Install YARN – On the New Node.

By now, you should have two nodes or servers:

Ensure to follow the naming conditions to avoid confusion. Its very important.

	<b>Container</b>	<b>Remarks</b>
<b>New Node</b>	Hadoopo	YARN Services
<b>Existing Node</b>	Sparko	Spark Client or Spark.

All the below commands should be executed on (Hadoop Node – **hadoopo**) unless specify. We are configuring YARN cluster now.

Change directory to the location where you have the software.

Extract the hadoop & jdk compressed files as shown below:

```
tar -xvf hadoop-* -C /opt  
tar -xvf jdk-11* -C /opt
```

Rename the folder for better readability:

```
#mv hadoo* hadoop  
#mv jdk* jdk
```

--- Install JDK and set Java Home (Copy the bin file in the YARN folder) , To include JAVA\_HOME for all bash users , make an entry in /etc/profile.d as follows:

```
echo "export JAVA_HOME=/opt/jdk/" > /etc/profile.d/java.sh
```

After Unpack the Hadoop distribution. In the distribution, edit the file `/opt/hadoop/etc/hadoop/hadoop-env.sh` to define some parameters as follows: Look for the below line and update.

```
# set to the root of your Java installation  
export JAVA_HOME=/opt/jdk  
  
# cd /opt/hadoop
```

Try the following command:

```
$ bin/hadoop
```

It should list commands as shown below, if everything is configured properly

```

kerbname      show auth_to_local principal conversion
key           manage keys via the KeyProvider
rumenfolder   scale a rumen input trace
rumentrace    convert logs into a rumen trace
s3guard       manage metadata on S3
trace         view and modify Hadoop tracing settings
version       print the version

```

Daemon Commands:

```
kms          run KMS, the Key Management Server
```

SUBCOMMAND may print help when invoked w/o parameters or with -h.

```
[root@hadoop0 hadoop]#
```

You need to modify some setting as follows: Replace with your hostname accordingly.

Use and Update the following:

Specify the Namenode hostname and its port.

```
#vi /opt/hadoop/etc/hadoop/core-site.xml
```

```

<configuration>
  <property>
    <name>fs.defaultFS</name>
    <value>hdfs://hadoop0:8020</value>

```

```
</property>
<property>
    <name>io.file.buffer.size</name>
    <value>131072</value>
    <description>Buffer size</description>
</property>
</configuration>
```

Configure its replication.

```
#vi /opt/hadoop/etc/hadoop/hdfs-site.xml
```

```
<configuration>
<property>
    <name>dfs.replication</name>
    <value>1</value>
</property>
</configuration>
```

Setup passphraseless ssh for the host to logon itself for hadoop services.

Now check that you can ssh to the **hadoop** without a passphrase:

```
yum -y install openssh-server openssh-clients
systemctl start sshd
```

Change the password of root using passwd command.

```
#passwd
```

Enter the password.

```
$ ssh hadoopo
```

If you cannot ssh to hadoopo without a passphrase, execute the following commands:

```
$ ssh-keygen -t rsa -P "" -f ~/.ssh/id_rsa
$ cat ~/.ssh/id_rsa.pub >> ~/.ssh/authorized_keys
$ chmod 0600 ~/.ssh/authorized_keys
```

Execute to initialize the required parameters on **hadoopo** terminal.

```
export HDFS_NAMENODE_USER="root"
export HDFS_DATANODE_USER="root"
export HDFS_SECONDARYNAMENODE_USER="root"
export YARN_RESOURCEMANAGER_USER="root"
export YARN_NODEMANAGER_USER="root"
```

Format the hdfs filesystem:

```
$ hdfs namenode -format
```

```
[root@hadoop0 hadoop]# 
2023-07-04 21:02:03,839 INFO common.Storage: Storage directory /opt/hdfs/namenode has been successfully formatted.
2023-07-04 21:02:04,017 INFO namenode.FSImageFormatProtobuf: Saving image file /opt/hdfs/namenode/current/fsimage.ckpt_0000000000
0000000000 using no compression
2023-07-04 21:02:04,565 INFO namenode.FSImageFormatProtobuf: Image file /opt/hdfs/namenode/current/fsimage.ckpt_000000000000000000
000 of size 391 bytes saved in 0 seconds .
2023-07-04 21:02:04,647 INFO namenode.NNStorageRetentionManager: Going to retain 1 images with txid >= 0
2023-07-04 21:02:04,696 INFO namenode.NameNode: SHUTDOWN_MSG:
*****
SHUTDOWN_MSG: Shutting down NameNode at hadoop0/172.18.0.2
*****/[root@hadoop0 hadoop]# 
```

Start NameNode daemon and DataNode daemon:

```
#cd /opt/hadoop
$ sbin/start-dfs.sh
```

```
[root@hadoop0 hadoop]# sbin/start-dfs.sh
Starting namenodes on [localhost]
Last login: Thu Jan 21 08:00:44 UTC 2021 from localhost on pts/2
Starting datanodes
Last login: Thu Jan 21 08:06:21 UTC 2021 on pts/2
Starting secondary namenodes [hadoop0]
Last login: Thu Jan 21 08:06:23 UTC 2021 on pts/2
hadoop0: Warning: Permanently added 'hadoop0,172.17.0.2' (ECDSA) to the list of known hosts.
[root@hadoop0 hadoop]# 
```

Browse the web interface for the NameNode; by default, it is available at:  
NameNode –

<http://localhost:9870/>

Verify the Node as shown below.



## Overview 'hadoop0:8020' (active)

Started:	Wed Jul 05 09:18:57 +0530 2023
Version:	3.1.2, r1019dde65bcf12e05ef48ac71e84550d589e5d9a
Compiled:	Tue Jan 29 07:09:00 +0530 2019 by sunilg from branch-3.1.2
Cluster ID:	CID-88e30be5-a9b0-4178-8048-61363734a0f7
Block Pool ID:	BP-1865112727-172.18.0.2-1688484723073

## Summary

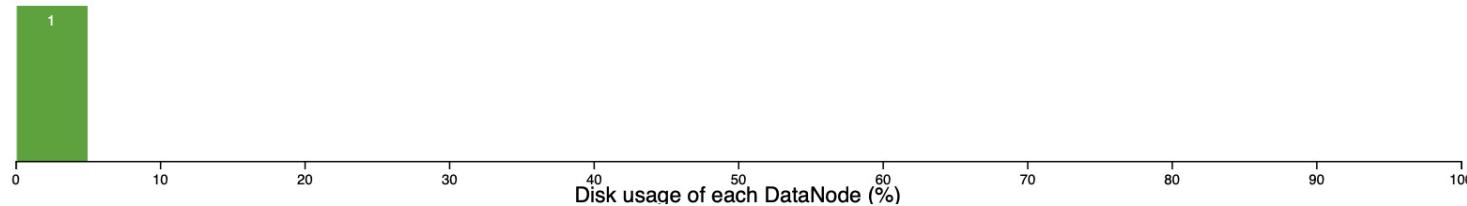
Click on Datanodes.



## Datanode Information

✓ In service    ⚠ Down    ⚡ Decommissioned    ⚡ Decommissioned & dead    🔍 In Maintenance    🔍 In Maintenance & dead

### Datanode usage histogram



### In operation

Show 25 entries		Search:					
Node	Http Address	Last contact	Last Block Report	Capacity	Blocks	Block pool used	Version
✓ hadoop0:9866 (172.18.0.2:9866)	<a href="http://hadoop0:9864">http://hadoop0:9864</a>	1s	1m	125.93 GB	0	24 KB (0%)	3.1.2
Showing 1 to 1 of 1 entries							
Previous	1	Next					

You should have one data node as shown above.

Make the HDFS directories required to execute MapReduce jobs:

```
$ hdfs dfs -mkdir /user  
$ hdfs dfs -mkdir /user/root
```

You can run a MapReduce job on YARN in a pseudo-distributed mode

Start HDFS and YARN on the YARN Node: **hadoop1**

Configure parameters as follows:

```
#vi /opt/hadoop/etc/hadoop/mapred-site.xml
```

```
<configuration>  
  <property>  
    <name>mapreduce.framework.name</name>  
    <value>yarn</value>  
  </property>  
  <property>  
    <name>mapreduce.application.classpath</name>  
    <value>$HADOOP_MAPRED_HOME/share/hadoop/mapreduce/*:$HADOOP_MAPRED_HOME/sh  
are/hadoop/mapreduce/lib/*</value>  
  </property>  
</configuration>
```

```
#vi /opt/hadoop/etc/hadoop/yarn-site.xml

<configuration>
<property>
    <name>yarn.resourcemanager.hostname</name>
    <value>hadoopo</value>
    <description>The hostname of the RM.</description>
</property>
<property>
    <name>yarn.nodemanager.aux-services</name>
    <value>mapreduce_shuffle</value>
</property>
<property>
    <name>yarn.nodemanager.env-whitelist</name>
<value>JAVA_HOME,HADOOP_COMMON_HOME,HADOOP_HDFS_HOME,HADOOP_CONF_DIR,
CLASSPATH_PREPEND_DISTCACHE,HADOOP_YARN_HOME,HADOOP_MAPRED_HOME</value>
>
</property>
<property>
<name>yarn.nodemanager.vmem-check-enabled</name>
<value>false</value>
<description>Whether virtual memory limits will be enforced for containers</description>
</property>
<property>
<name>yarn.nodemanager.vmem-pmem-ratio</name>
<value>4</value>
<description>Ratio between virtual memory to physical memory when setting memory limits
for containers</description>
</property>
```

```
</configuration>
```

Start ResourceManager daemon and NodeManager daemon

Open a terminal to hadoop

```
#cd /opt/hadoop  
$ sbin/start-yarn.sh
```

```
[root@hadoop0 hadoop]# sbin/start-yarn.sh  
Starting resourcemanager  
Last login: Thu Jan 21 08:06:32 UTC 2021 on pts/2  
Starting nodemanagers  
Last login: Thu Jan 21 08:13:09 UTC 2021 on pts/2  
[root@hadoop0 hadoop]#
```

Browse the web interface for the ResourceManager; by default it is available at:

ResourceManager - <http://localhost:8088/>

#verify with jps , all the below services should be there.

```
jps
```

```
[root@hadoop0 hadoop]# export PATH=$PATH:/opt/jdk/bin
[root@hadoop0 hadoop]# jps
1826 NodeManager
919 NameNode
1034 DataNode
1210 SecondaryNameNode
1708 ResourceManager
2174 Jps
[root@hadoop0 hadoop]#
```

Hadoop Node - Let us set some environment and path variable as follows: ([vi ~/.bashrc](#))

```
export HADOOP_HOME=/opt/hadoop
export PATH=$HADOOP_HOME/bin:$PATH
```

Let us create a temporary folder, that will be used for working space for the YARN cluster.

```
hdfs dfs -mkdir /tmp
hdfs dfs -chmod -R 1777 /tmp
hdfs dfs -mkdir /tmp/in
hdfs dfs -ls /tmp
```

You can get the README.md file from the Software folder, which is provided along with the training.

```
hdfs dfs -copyFromLocal /opt/hadoop/README.txt /tmp/in
hdfs dfs -ls /tmp/in
```

[hdfs dfs -mkdir /tmp/spark-events](#)

Congrats! You have successfully configured Yarn Cluster.

Let us configure hadoop client setting on the Spark node -> **sparko**.

Compress the Hadoop folder on the Hadoop instances or Hadoop node.

```
#cd /opt  
#tar -cvzf hadoop.tar hadoop/
```

Copy the compressed hadoop folder to the sparko or spark node and uncompressed in /opt folder.

[Docker:

```
copy from hadoop to host : docker cp hadoop:/opt/hadoop.tar .  
copy from host to spark : docker cp hadoop.tar sparko:/opt/  
docker cp hadoop:/opt/hadoop.tar sparko:/opt/  
docker cp hadoop.tar hadoop1:/opt/hadoop.tar  
]
```

On the **sparko** or spark instance.

```
#tar -xvf hadoop.tar
```

Update **yarn-site.xml** on the Spark Node.

```
#vi /opt/spark/etc/hadoop/yarn-site.xml
```

```
<property>  
  <name>yarn.resourcemanager.hostname</name>  
  <value>hadoopo</value>  
  <description>The hostname of the RM.</description>  
</property>  
<property>
```

```
<name>yarn.resourcemanager.address</name>
<value>hadoop:8032</value>
<description>The hostname of the RM.</description>
</property>
```

Create the following folder on HDFS cluster and copy the file in the folder.

```
# hdfs dfs -mkdir -p /opt/spark
# hdfs dfs -copyFromLocal /opt/spark/README.md /opt/spark/
```

You can verify the file with the following command. This command will display the content of the file in the HDFS cluster.

```
# hdfs dfs -cat /opt/spark/README.md
```

Open a command console to execute the Spark jobs to submit on YARN.

Ensure that you install python on Hadoop node:

```
# yum install -y python3
```

On Spark Node

Export Hadoop and Yarn environment variable pointing to the Hadoop Node – Hadoop conf file.

```
export JAVA_HOME=/opt/jdk
export HADOOP_HOME=/opt/hadoop
export PATH=$PATH:$JAVA_HOME/bin:$HADOOP_HOME/sbin:$HADOOP_HOME/bin
```

Create a python app, **SimpleApp.py** with the following code:

```
# vi /opt/SimpleApp.py
// ----- Code Begin -----//

"""SimpleApp.py"""
from pyspark.sql import SparkSession

logFile = "/opt/spark/README.md" # Should be some file on your system

spark = SparkSession.builder.appName("Python App").getOrCreate()
logData = spark.read.text(logFile).cache()

numAs = logData.filter(logData.value.contains('a')).count()
numBs = logData.filter(logData.value.contains('b')).count()

print("Lines with a: %i, lines with b: %i" % (numAs, numBs))

spark.stop()
```

```
// ----- Code Ends -----//
```

The above code reads the README.md file from HDFS cluster and counts the line containing 'a' and 'b' respectively.

On **spark0** terminal

```
#export HADOOP_CONF_DIR=/opt/hadoop/etc/hadoop
#export YARN_CONF_DIR=/opt/hadoop/etc/hadoop
#cd /opt/spark/bin
./spark-submit --master yarn --deploy-mode cluster --executor-memory 512m --num-executors 2
/opt/SimpleApp.py
```

or client side - Optional.

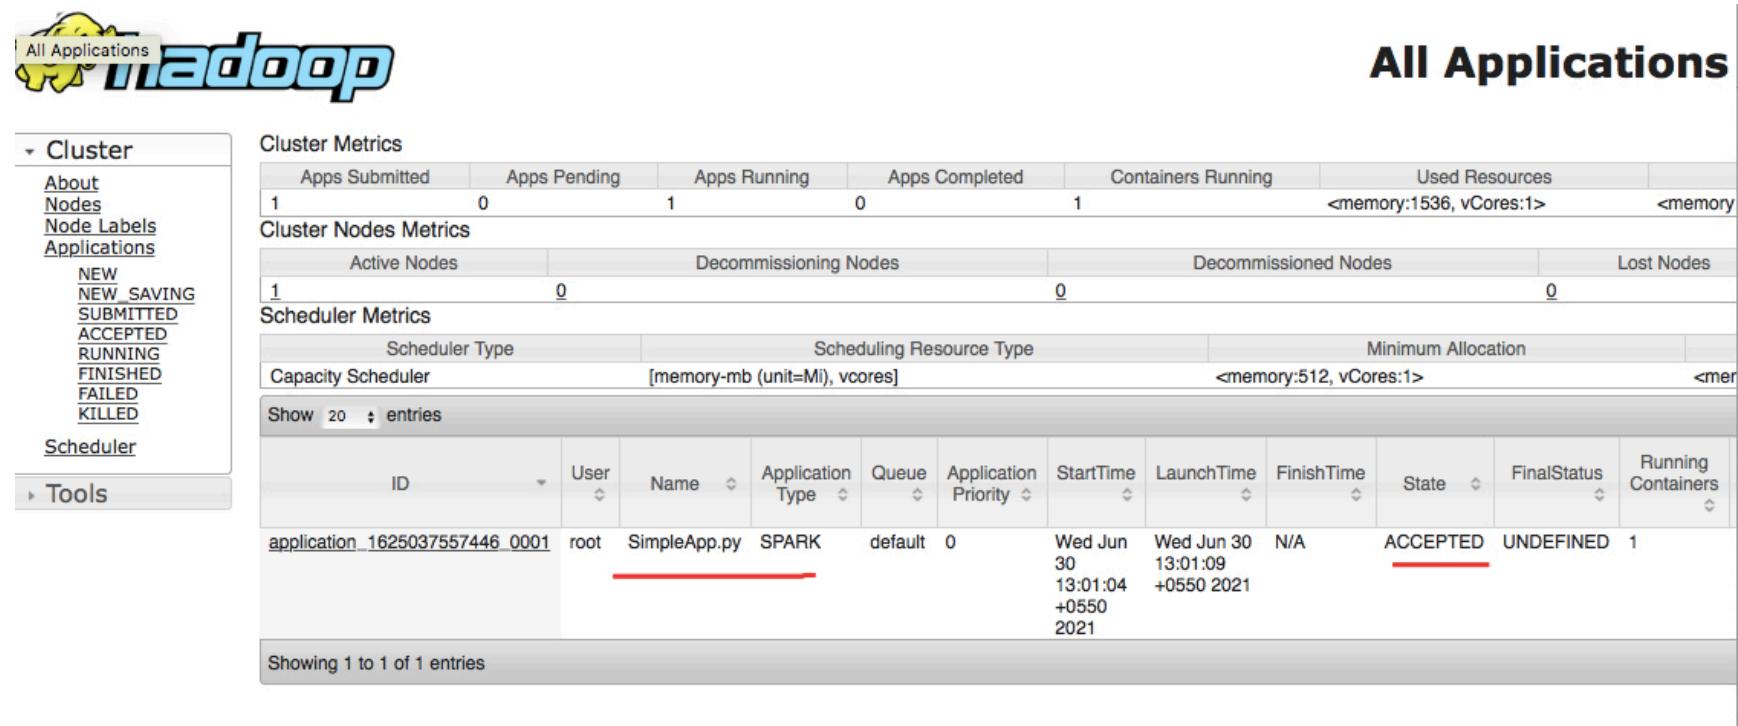
```
#./spark-submit --master yarn --deploy-mode client /opt/SimpleAppP.py
```

Ensure that `HADOOP_CONF_DIR` or `YARN_CONF_DIR` points to the directory which contains the (client side) configuration files for the Hadoop cluster.

Note --> Application py/Jar should be in the local file system, In folder is in HDFS and Out should not be present in the HDFS,it will be automatically created.

You can verify the progress of the application using <http://localhost:8088/cluster>

Click on Cluster-->Application --> accepted



The screenshot shows the Apache Hadoop 'All Applications' interface. On the left, there's a sidebar with navigation links: 'All Applications', 'Cluster Metrics', 'About Nodes', 'Node Labels', 'Applications' (which is currently selected), 'Scheduler', and 'Tools'. The main area displays 'Cluster Metrics' with values: Apps Submitted (1), Apps Pending (0), Apps Running (1), Apps Completed (0), Containers Running (1), Used Resources (<memory:1536, vCores:1>), and Lost Nodes (<memory>). Below that is 'Cluster Nodes Metrics' showing Active Nodes (1), Decommissioning Nodes (0), Decommissioned Nodes (0), and Lost Nodes (0). Under 'Scheduler Metrics', it shows Scheduler Type (Capacity Scheduler) and Scheduling Resource Type ([memory-mb (unit=Mi), vcores]). The table lists one application entry:

ID	User	Name	Application Type	Queue	Application Priority	StartTime	LaunchTime	FinishTime	State	FinalStatus	Running Containers
application_1625037557446_0001	root	SimpleApp.py	SPARK	default	0	Wed Jun 30 13:01:04 2021	13:01:09 +0550	N/A	ACCEPTED	UNDEFINED	1

At the bottom, a message says 'Showing 1 to 1 of 1 entries'.

Wait for few moment and verify on Running tab.

ID	User	Name	Application Type	Queue	Application Priority	StartTime	LaunchTime	FinishTime	State	FinalStatus	Running Containers
application_1625129460519_0002	root	SimpleApp.py	SPARK	default	0	Thu Jul 1 14:36:52 +0550 2021	Thu Jul 1 14:36:52 +0550 2021	N/A	RUNNING	UNDEFINED	3

Click on the Application ID --> Logs

The screenshot shows the Hadoop Cluster Application Overview page. The URL is [http://hp.com:8088/cluster/app/application\\_1434297324587\\_0003](http://hp.com:8088/cluster/app/application_1434297324587_0003). The page displays the following information:

**Application Overview**

- User: root
- Name: com.ht.sparks.WordCount
- Application Type: SPARK
- Application Tags:
- State: RUNNING
- FinalStatus: UNDEFINED
- Started: 14-Jun-2015 23:25:58
- Elapsed: 3mins, 57sec
- Tracking URL: [ApplicationMaster](#)
- Diagnostics:

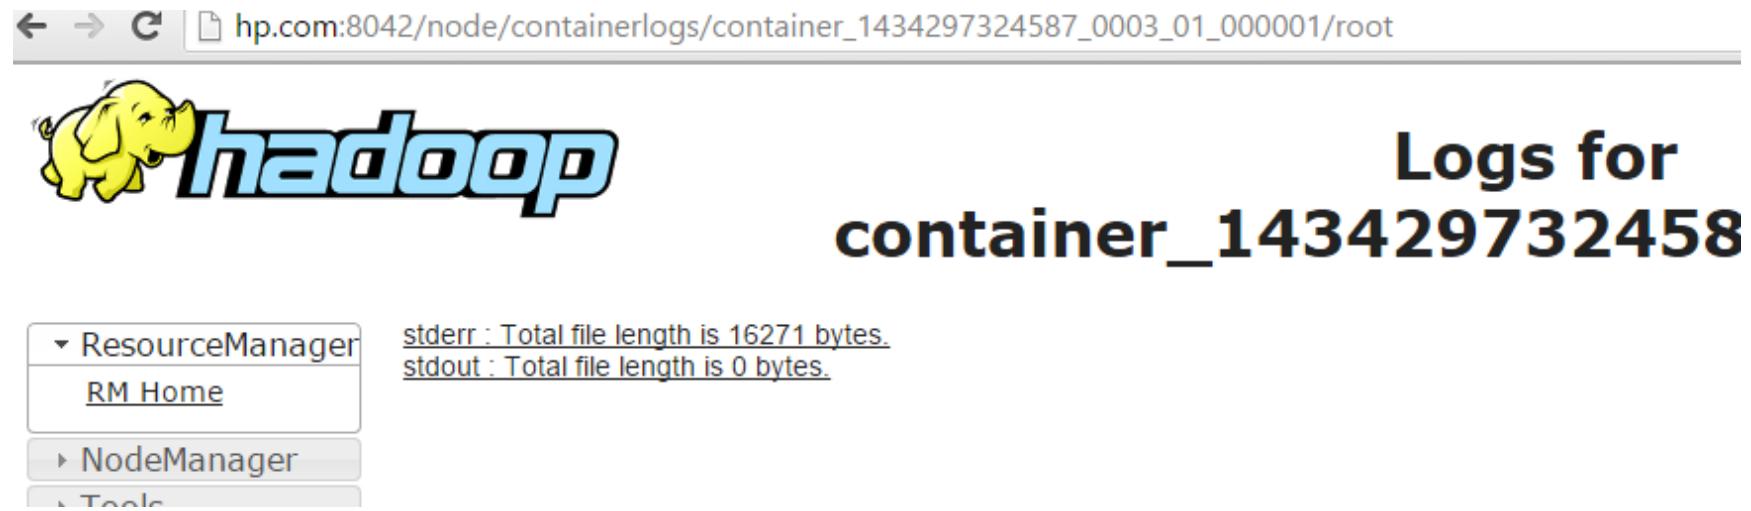
**Application Metrics**

- Total Resource Preempted: <memory:0, vCores:0>
- Total Number of Non-AM Containers Preempted: 0
- Total Number of AM Containers Preempted: 0
- Resource Preempted from Current Attempt: <memory:0, vCores:0>
- Number of Non-AM Containers Preempted from Current Attempt: 0
- Aggregate Resource Allocation: 701453 MB-seconds, 459 vcore-seconds

**ApplicationMaster**

Attempt Number	Start Time	Node	Logs
1	14-Jun-2015 23:25:58	<a href="#">hp.com:8042</a>	<a href="#">logs</a>

Logs --> Stderr to get details as follows:



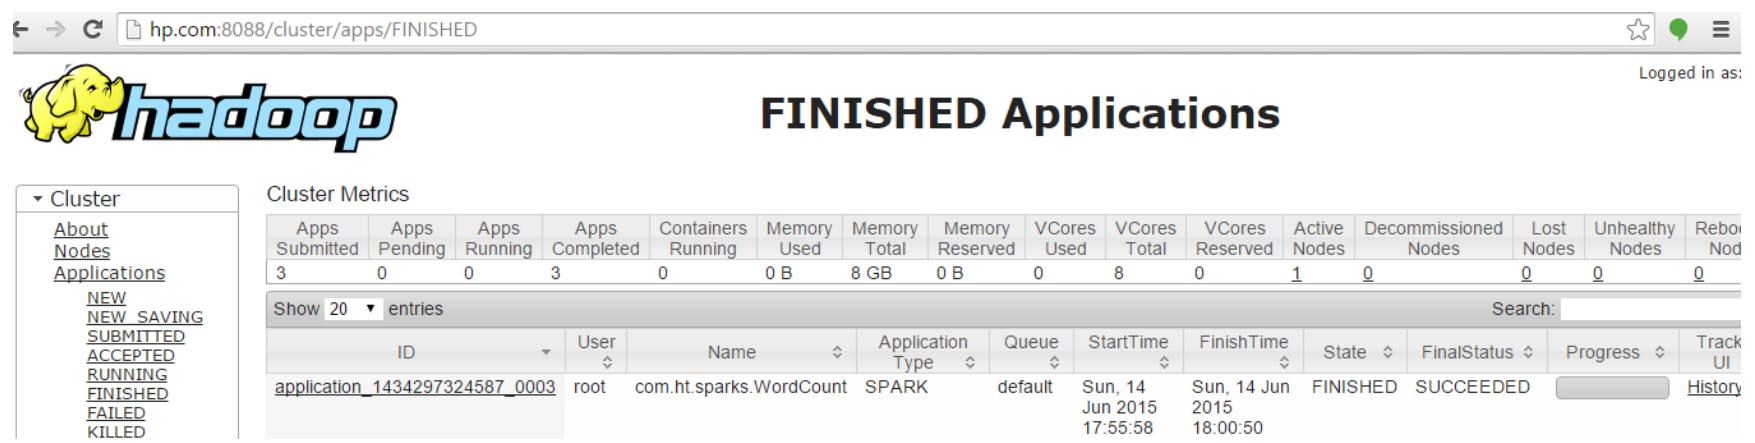
The screenshot shows the Hadoop ResourceManager UI at [http://hp.com:8042/node/containerlogs/container\\_1434297324587\\_0003\\_01\\_000001/root](http://hp.com:8042/node/containerlogs/container_1434297324587_0003_01_000001/root). The page title is "Logs for container\_1434297324587". On the left, there's a sidebar with links for ResourceManager (RM Home), NodeManager, and Tools. The main content area displays log entries:

```

stderr : Total file length is 16271 bytes.
stdout : Total file length is 0 bytes.

```

**After sometimes click on Finished, after the console exit the program.**



The screenshot shows the Hadoop Cluster UI at <http://hp.com:8088/cluster/apps/FINISHED>. The page title is "FINISHED Applications". On the left, there's a sidebar with a "Cluster" section containing links for About, Nodes, Applications, and a status table:

NEW	NEW_SAVING	SUBMITTED	ACCEPTED	RUNNING	FINISHED	FAILED	KILLED
-----	------------	-----------	----------	---------	----------	--------	--------

The main content area shows "Cluster Metrics" and a table of finished applications:

Apps Submitted	Apps Pending	Apps Running	Apps Completed	Containers Running	Memory Used	Memory Total	Memory Reserved	VCores Used	VCores Total	VCores Reserved	Active Nodes	Decommissioned Nodes	Lost Nodes	Unhealthy Nodes	Reboot Nod
3	0	0	3	0	0 B	8 GB	0 B	0	8	0	1	0	0	0	0

Below this is a table of finished applications:

ID	User	Name	Application Type	Queue	StartTime	FinishTime	State	FinalStatus	Progress	Track UI
application_1434297324587_0003	root	com.ht.sparks.WordCount	SPARK	default	Sun, 14 Jun 2015 17:55:58	Sun, 14 Jun 2015 18:00:50	FINISHED	SUCCEEDED		<a href="#">History</a>

On the job submit console.

```
2021-01-22 16:36:11,589 INFO yarn.Client: Application report for application_1611331123305_0002 (state: RUNNING)
2021-01-22 16:36:12,662 INFO yarn.Client: Application report for application_1611331123305_0002 (state: RUNNING)
2021-01-22 16:36:13,765 INFO yarn.Client: Application report for application_1611331123305_0002 (state: RUNNING)
2021-01-22 16:36:14,896 INFO yarn.Client: Application report for application_1611331123305_0002 (state: RUNNING)
2021-01-22 16:36:16,011 INFO yarn.Client: Application report for application_1611331123305_0002 (state: RUNNING)
2021-01-22 16:36:17,296 INFO yarn.Client: Application report for application_1611331123305_0002 (state: RUNNING)
2021-01-22 16:36:18,725 INFO yarn.Client: Application report for application_1611331123305_0002 (state: RUNNING)
2021-01-22 16:36:19,914 INFO yarn.Client: Application report for application_1611331123305_0002 (state: RUNNING)
2021-01-22 16:36:21,041 INFO yarn.Client: Application report for application_1611331123305_0002 (state: FINISHED)
2021-01-22 16:36:21,401 INFO yarn.Client:
    client token: N/A
    diagnostics: N/A
    ApplicationMaster host: hadoop0
    ApplicationMaster RPC port: 38603
    queue: default
    start time: 1611333080086
    final status: SUCCEEDED
    tracking URL: http://hadoop0:8088/proxy/application_1611331123305_0002/
    user: root
2021-01-22 16:36:23,265 INFO util.ShutdownHookManager: Shutdown hook called
2021-01-22 16:36:23,442 INFO util.ShutdownHookManager: Deleting directory /tmp/spark-7b20b446-2b75-4550-9e2d-c84ad6
7c9b3b
2021-01-22 16:36:23,532 INFO util.ShutdownHookManager: Deleting directory /tmp/spark-89770943-febc-4bdb-bf6e-404103
88026d
[root@303b68143644 bin]#
```

Since we have printed the output. You can verify from the log file only.

Go to the User Log folders of the hadoop installation. (Job ID and the container ID will be different in your machine, update accordingly)

```
#cd /opt/hadoop/logs/userlogs  
#more application_1625129460519_0004/container_1625129460519_0004_01_000001/stdout
```

```
[root@hadoop0 userlogs]# pwd  
/opt/hadoop/logs/userlogs  
[root@hadoop0 userlogs]# more application_1625129460519_0004/container_1625129460519_0004_01_000001/stdout  
Lines with a: 64, lines with b: 32  
[root@hadoop0 userlogs]#
```

You can also view the log from Web UI

Directory → Utilities -> Logs → userlogs -> application id (Your application id) -> Container Id(browse subsequently all the logs) --> stdout.

```
23/07/07 10:16:49 INFO YarnClusterScheduler: Killing all running tasks in stage 5: Stage finished  
23/07/07 10:16:49 INFO DAGScheduler: Job 3 finished: count at NativeMethodAccessorImpl.java:0, took 7.896622 s  
Lines with a: 71, lines with b: 38  
23/07/07 10:16:49 WARN AbstractConnector:  
java.io.IOException: No such file or directory
```

Task -> Let us write the output into file instead of writing in console.

Create a file **SimpleAppP.py** and update with the following code. It only writes the dataframe which has the line having 'a' in it.

//----- Code Begin -----

```
"""SimpleApp.py"""
from pyspark.sql import SparkSession

logFile = "/opt/spark/README.md" # Should be some file on your system

spark = SparkSession.builder.appName("Python App").getOrCreate()
logData = spark.read.text(logFile).cache()

numAs = logData.filter(logData.value.contains('a')).count()
numBs = logData.filter(logData.value.contains('b')).count()

#print("Lines with a: %i, lines with b: %i" % (numAs, numBs))

myDF = logData.filter(logData.value.contains('a'))
myDF.write.save("mydatap")
spark.stop()
```

//----- Code Ends -----

Execute the program again.

```
#spark-submit --master yarn --deploy-mode cluster --executor-memory 512m --num-executors 2
SimpleAppP.py
```

At the end of the execution, you should have the following output in the console.

```
2021-07-01 10:10:44,466 INFO yarn.Client: Application report for application_1625129460519_0005 (state: FINISHED)
2021-07-01 10:10:44,467 INFO yarn.Client:
  client token: N/A
  diagnostics: N/A
  ApplicationMaster host: hadoop0
  ApplicationMaster RPC port: 41043
  queue: default
  start time: 1625133769650
  final status: SUCCEEDED
  tracking URL: http://hadoop0:8088/proxy/application_1625129460519_0005/
  user: root
2021-07-01 10:10:44,535 INFO util.ShutdownHookManager: Shutdown hook called
2021-07-01 10:10:44,550 INFO util.ShutdownHookManager: Deleting directory /tmp/spark-85f5b09e-f56b-49a9-9ca7-fad964848b7f
2021-07-01 10:10:44,835 INFO util.ShutdownHookManager: Deleting directory /tmp/spark-009ed528-fe7b-4f7e-acda-14c34a265a3a
[root@spark0 data]#
```

You can verify the output file in the hdfs as shown below:

```
#hdfs dfs -ls -R /user/root/mydatap
```

```
23/07/07 12:26:35 INFO ShutdownHookManager: Deleting directory /tmp/spark-fa9f16f5-b4a8-405d-a39f-1079a24a109e
23/07/07 12:26:35 INFO ShutdownHookManager: Deleting directory /tmp/spark-4183b6bb-c10a-4bc7-9553-cfd6d87ac93d
[root@hadoop1 opt]# hdfs dfs -ls -R /user/root/mydatap
-rw-r--r-- 2 root supergroup 0 2023-07-07 12:25 /user/root/mydatap/_SUCCESS
-rw-r--r-- 2 root supergroup 3464 2023-07-07 12:25 /user/root/mydatap/part-00000-713d9b32-24cb-47ed-b6e1-7044ca7ccaf9-c000.snappy.parquet
[root@hadoop1 opt]#
```

Substitute the path of the file listed by above command.

```
# hdfs dfs -cat /user/root/mydatap/part-00000-713d9b32-24cb-47ed-b6e1-7044ca7ccaf9-c000.snappy.parquet
```

to AlkaHDFSeother  
of ch different vocs  
=M r must% g t\sa%  
k-aUA ar#skguiU&eV"aa ticular distribu  
a A2?pa=Hm  
(Thriftserve6s.I ConfiguI bZ[2& G  
np=Hc.`B  
Fok  
6 \$n overview  
m. [ . A Ato n /  
)wng &g

It's a binary parquet file.

Great! You have successfully executed spark job in YARN cluster.

Let us verify the output now using HDFS browser.

<http://localhost:9870/explorer.html#/>

Click Utilities --> Browse the file system → /user/root/output (Enter in the Directory ) - Go

The screenshot shows the Hadoop Web UI interface. At the top, there is a green header bar with several tabs: Hadoop, Overview, Datanodes, Datanode Volume Failures, Snapshot, Startup Progress, and Utilities. The Utilities tab is highlighted with a red circle and has a red arrow pointing to it from the text above. Below the header, the page title is "Browse Directory". In the search bar, the path "/user/root/mydatap" is entered, with a "Go!" button and three icons (file, refresh, and list) to its right. Underneath, there are filters for "Show 25 entries" and a "Search:" input field. The main content is a table listing two files in the directory:

	Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
<input type="checkbox"/>	-rw-r--r--	root	supergroup	0 B	Jul 01 15:40	2	128 MB	_SUCCESS
<input type="checkbox"/>	-rw-r--r--	root	supergroup	3.34 KB	Jul 01 15:40	2	128 MB	part-00000-554ff738-ee36-449c-b654-9a7fa92f17b0-c000.snappy.parquet

At the bottom left, it says "Showing 1 to 2 of 2 entries". On the right, there are "Previous" and "Next" buttons, with the number "1" in the center, all enclosed in a red horizontal bar. A red arrow also points from the "Utilities" label to the "part-00000..." file name in the table.

Click on the above mark file.

## Browse Directory

/tmp/out							Go!
Permission	Owner	Group	Size	Replication	Block Size	Name	
-rw-r--r--	root	supergroup	0 B	3	128 MB	_SUCCESS	
-rw-r--r--	root	supergroup	60.71 KB	3	128 MB	part-00000	
-rw-r--r--	root	supergroup	61.17 KB	3	128 MB	part-00001	
-rw-r--r--	root	supergroup	60.76 KB	3	128 MB	part-00002	
-rw-r--r--	root	supergroup	57.2 KB	3	128 MB	part-00003	

Let us view one of the output. click on part-0001 --> Download.

When you're done, you can stop the daemons with: Optional.  
On the hadoop terminal

```
# stop-yarn.sh  
#stop-hdfs.sh
```

Congrats!.

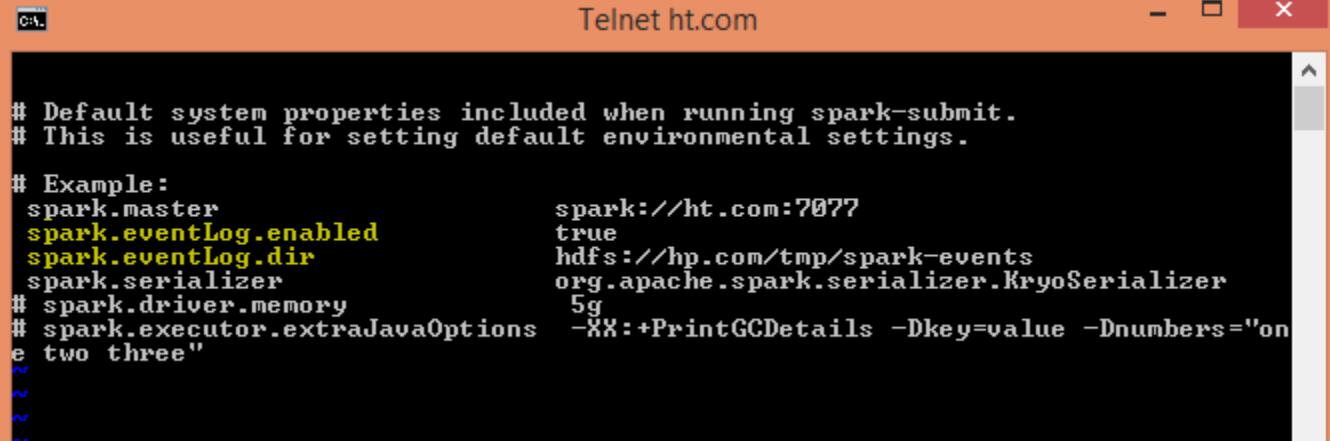
**----- End of Lab-----**

**Errata:**

- 1) 15/06/14 21:50:46 INFO storage.BlockManagerMasterActor: Registering block manager ht.com:50475 with 267.3 MB RAM, BlockManagerId(<driver>, ht.com, 50475)  
15/06/14 21:50:46 INFO storage.BlockManagerMaster: Registered BlockManager  
Exception in thread "main" java.lang.IllegalArgumentException: Wrong FS: file:/tmp/spark-events/application\_1434297324587\_0001.inprogress, expected: hdfs://hp.com at org.apache.hadoop.fs.FileSystem.checkPath(FileSystem.java:643)  
at org.apache.hadoop.hdfs.DistributedFileSystem.getTableName(DistributedFileSystem.java:191)  
at org.apache.hadoop.hdfs.DistributedFileSystem.access\$000(DistributedFileSystem.java:102)  
at org.apache.hadoop.hdfs.DistributedFileSystem\$22.doCall(DistributedFileSystem.java:1266)  
at org.apache.hadoop.hdfs.DistributedFileSystem\$22.doCall(DistributedFileSystem.java:1262)  
at org.apache.hadoop.fs.FileSystemLinkResolver.resolve(FileSystemLinkResolver.java:81)  
at org.apache.hadoop.hdfs.DistributedFileSystem.setPermission(DistributedFileSystem.java:1262)  
at org.apache.spark.scheduler.EventLoggingListener.start(EventLoggingListener.scala:128)

```
15/06/14 21:50:46 INFO storage.BlockManagerMasterActor: Registering block manager ht.com:50475 with 267.3 MB RAM, BlockManagerId(<driver>, ht.com, 50475)
15/06/14 21:50:46 INFO storage.BlockManagerMaster: Registered BlockManager
Exception in thread "main" java.lang.IllegalArgumentException: Wrong FS: file:/tmp/spark-events/application_1434297324587_0001.inprogress, expected: hdfs://hp.com
    at org.apache.hadoop.fs.FileSystem.checkPath(FileSystem.java:643)
    at org.apache.hadoop.hdfs.DistributedFileSystem.getTableName(DistributedFileSystem.java:191)
    at org.apache.hadoop.hdfs.DistributedFileSystem.access$000(DistributedFileSystem.java:102)
    at org.apache.hadoop.hdfs.DistributedFileSystem$22.doCall(DistributedFileSystem.java:1266)
    at org.apache.hadoop.hdfs.DistributedFileSystem$22.doCall(DistributedFileSystem.java:1262)
    at org.apache.hadoop.fs.FileSystemLinkResolver.resolve(FileSystemLinkResolver.java:81)
    at org.apache.hadoop.hdfs.DistributedFileSystem.setPermission(DistributedFileSystem.java:1262)
    at org.apache.spark.scheduler.EventLoggingListener.start(EventLoggingListener.scala:128)
    at org.apache.spark.SparkContext.<init>(SparkContext.scala:399)
    at org.apache.spark.api.java.JavaSparkContext.<init>(JavaSparkContext.sc
ala:61)
```

Solution: Verify default configuration of the spark installation. /spark/spark-1.3.0-bin-hadoop2.4/conf/spark-defaults.conf



```
# Default system properties included when running spark-submit.  
# This is useful for setting default environmental settings.  
  
# Example:  
spark.master          spark://ht.com:7077  
spark.eventLog.enabled true  
spark.eventLog.dir    hdfs://ht.com/tmp/spark-events  
spark.serializer      org.apache.spark.serializer.KryoSerializer  
# spark.driver.memory  5g  
# spark.executor.extraJavaOptions -XX:+PrintGCDetails -Dkey=value -Dnumbers="one  
two three"  
  
~  
~
```

Check that spark.eventLog.dir is begin with hdfs:// and the /tmp/spark-events is already created in HDFS system

You can verify on the YARN cluster only.

Verify using : hadoop fs -ls -R /tmp/

```
[root@hp ~]# hadoop fs -ls -R /tmp/
Java HotSpot(TM) Client VM warning: You have loaded library /YARN/hadoop-2.6.0/lib/native/libhadoop.so.1.0.0 which might have disabled stack guard. The VM will try to fix the stack guard now.
It's highly recommended that you fix the library with 'execstack -c <libfile>', or link it with '-z noexecstack'.
15/06/14 22:04:11 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
drwxr-xr-x  - root supergroup          0 2015-06-14 16:25 /tmp/in
-rw-r--r--  1 root supergroup 3629 2015-06-14 16:25 /tmp/in/README.md
drwxr-xr-x  - root supergroup          0 2015-06-14 17:47 /tmp/out
-rw-r--r--  3 root supergroup          0 2015-06-14 17:47 /tmp/out/_SUCCESS
-rw-r--r--  3 root supergroup 1956 2015-06-14 17:47 /tmp/out/part-00000
-rw-r--r--  3 root supergroup 1717 2015-06-14 17:47 /tmp/out/part-00001
drwxr-xr-x  - root supergroup          0 2015-06-14 17:50 /tmp/out1
-rw-r--r--  3 root supergroup          0 2015-06-14 17:50 /tmp/out1/_SUCCESS
-rw-r--r--  3 root supergroup 1956 2015-06-14 17:50 /tmp/out1/part-00000
-rw-r--r--  3 root supergroup 1717 2015-06-14 17:50 /tmp/out1/part-00001
drwxr-xr-x  - root supergroup          0 2015-06-14 17:52 /tmp/out3
-rw-r--r--  3 root supergroup          0 2015-06-14 17:52 /tmp/out3/_SUCCESS
-rw-r--r--  3 root supergroup 1956 2015-06-14 17:52 /tmp/out3/part-00000
-rw-r--r--  3 root supergroup 1717 2015-06-14 17:52 /tmp/out3/part-00001
drwxr-xr-x  - root supergroup          0 2015-06-14 21:59 /tmp/out5
-rw-r--r--  3 root supergroup          0 2015-06-14 21:59 /tmp/out5/_SUCCESS
-rw-r--r--  3 root supergroup 1956 2015-06-14 21:59 /tmp/out5/part-00000
-rw-r--r--  3 root supergroup 1717 2015-06-14 21:59 /tmp/out5/part-00001
drwxr-xr-x  - root supergroup          0 2015-06-14 21:58 /tmp/spark-events
-rw-rw-r--  3 root supergroup 20287 2015-06-14 21:59 /tmp/spark-events/app
lication_1434297324587_0002.inprogress
[...]
```

2) 15/06/14 00:41:20 INFO client.RMProxy: Connecting to ResourceManager at /o.o.o.o:8032

15/06/14 21:58:27 INFO yarn.Client: Application report for application\_1434297324587\_0002 (state: ACCEPTED)

Indefinite loop of above

**Solution:** Ensure that all configuration file are map to appropriate hostname and hostname to ip details are modified in the hosts file , Increase the resources and wait for few minutes to convert the status to running.

```
Yarn-site.xml
yarn.scheduler.minimum-allocation-mb: 256m
yarn.scheduler.increment-allocation-mb: 256m
```

```
<property>
<name>yarn.nodemanager.resource.memory-mb</name>
<value>4096</value> <!-- 4 GB -->
</property>
<property>
<name>yarn.scheduler.minimum-allocation-mb</name>
<value>512</value> <!-- 1 GB -->
</property>
<property>
<name>yarn.scheduler.increment-allocation-mb</name>
<value>256</value> <!-- 1 GB -->
</property>
```

**yarn-site.xml** – Spark Node. (Specify the hostname of the Hadoop RM Node as shown below)

```
<property>
  <name>yarn.resourcemanager.hostname</name>
  <value>hadoopo</value>
  <description>The hostname of the RM.</description>
</property>
<property>
  <name>yarn.resourcemanager.address</name>
  <value>hadoopo:8032</value>
  <description>The hostname of the RM.</description>
</property>
```

3) If unhealthy nodes are being displayed in the cluster URL with 1/1 local-dirs are bad: /tmp/hadoop-yarn/nm-local-dir;

**Solution:** delete the folder using the following command and ensure that nodes are healthy before proceeding forward, you can restart if required

```
hadoop fs -rm -fr /tmp/hadoop-yarn/*
```

Try the following command:

```
yarn application -list  
yarn application -kill [application name]
```

You can view jps in YARN cluster when spark job is executing. It will start Coarse\* processes.

```
[root@hp ~]# jps  
5082 DataNode  
9633 CoarseGrainedExecutorBackend  
9670 Jps  
4980 NameNode  
5345 NodeManager  
9625 CoarseGrainedExecutorBackend  
5029 SecondaryNameNode  
9420 ApplicationMaster  
5294 ResourceManager  
9629 CoarseGrainedExecutorBackend  
[root@hp ~]#
```

Even, cluster mode works

```
15/06/14 23:43:46 INFO yarn.Client: Application report for application_143429732
4587_0004 (state: FINISHED)
15/06/14 23:43:46 INFO yarn.Client:
  client token: N/A
  diagnostics: N/A
  ApplicationMaster host: hp.com
  ApplicationMaster RPC port: 0
  queue: default
  start time: 1434305512588
  final status: SUCCEEDED
  tracking URL: http://hp.com:8088/proxy/application_1434297324587_0004/
  user: root
[root@ht spark-1.3.0-bin-hadoop2.4]# ./bin/spark-submit --master yarn-cluster --
--class com.ht.sparks.WordCount --num-executors 3 --executor-cores 1 /spark/Sp
arkWC-0.0.1-SNAPSHOT.jar /tmp/in /tmp/out?
```

Issue:

```
UpToMaximumCountWithFixedSleep(maxRetries=10, sleepTime=1000 MILLISECONDS)
21/01/23 05:38:08 INFO Client: Retrying connect to server: 0.0.0.0/0.0.0.0:8032. Already tried 2 time(s); retry policy is Retry
UpToMaximumCountWithFixedSleep(maxRetries=10, sleepTime=1000 MILLISECONDS)
21/01/23 05:38:09 INFO Client: Retrying connect to server: 0.0.0.0/0.0.0.0:8032. Already tried 3 time(s); retry policy is Retry
UpToMaximumCountWithFixedSleep(maxRetries=10, sleepTime=1000 MILLISECONDS)
21/01/23 05:38:10 INFO Client: Retrying connect to server: 0.0.0.0/0.0.0.0:8032. Already tried 4 time(s); retry policy is Retry
UpToMaximumCountWithFixedSleep(maxRetries=10, sleepTime=1000 MILLISECONDS)
21/01/23 05:38:11 INFO Client: Retrying connect to server: 0.0.0.0/0.0.0.0:8032. Already tried 5 time(s); retry policy is Retry
UpToMaximumCountWithFixedSleep(maxRetries=10, sleepTime=1000 MILLISECONDS)
21/01/23 05:38:12 INFO Client: Retrying connect to server: 0.0.0.0/0.0.0.0:8032. Already tried 6 time(s); retry policy is Retry
UpToMaximumCountWithFixedSleep(maxRetries=10, sleepTime=1000 MILLISECONDS)
21/01/23 05:38:13 INFO Client: Retrying connect to server: 0.0.0.0/0.0.0.0:8032. Already tried 7 time(s); retry policy is Retry
UpToMaximumCountWithFixedSleep(maxRetries=10, sleepTime=1000 MILLISECONDS)
21/01/23 05:38:14 INFO Client: Retrying connect to server: 0.0.0.0/0.0.0.0:8032. Already tried 8 time(s); retry policy is Retry
UpToMaximumCountWithFixedSleep(maxRetries=10, sleepTime=1000 MILLISECONDS)
21/01/23 05:38:15 INFO Client: Retrying connect to server: 0.0.0.0/0.0.0.0:8032. Already tried 9 time(s); retry policy is Retry
UpToMaximumCountWithFixedSleep(maxRetries=10, sleepTime=1000 MILLISECONDS)
21/01/23 05:38:15 INFO RetryInvocationHandler: java.net.ConnectException: Your endpoint configuration is wrong; For more details see: http://wiki.apache.org/hadoop/UnsetHostnameOrPort, while invoking ApplicationClientProtocolPBClientImpl.getClusterMetrics over null after 1 failover attempts. Trying to failover after sleeping for 42760ms.
```

Verify the ip of the resource manager and set the home directory appropriately.  
Try updating the following setting in yarn-site.xml of Hadoop and spark node too..

```
<property>
  <name>yarn.resourcemanager.hostname</name>
  <value>hadoopo</value>
  <description>The hostname of the RM.</description>
</property>
<property>
  <name>yarn.resourcemanager.address</name>
  <value>hadoopo:8032</value>
  <description>The hostname of the RM.</description>
</property>
```

---

## 17. Jobs Monitoring : Using Web UI – 45 Minutes

Execute the Job as specified in the **Spark Two Nodes** cluster before proceeding ahead.

- Start the Two Nodes Spark Cluster
- Copy the input file in both the nodes, if you are executing on the Standalone cluster.

Execute the following job:

Update File, **/Software/data/names.json** with the following content:

```
{"firstName":"Grace","lastName":"Hopper"}  
 {"firstName":"Alan","lastName":"Turing"}  
 {"firstName":"Ada","lastName":"Lovelace"}  
 {"firstName":"Charles","lastName":"Babbage"}
```

Enter the following commands one at a time.

```
#pyspark --master spark://sparko:8090
```

```
[root@e5217f917ab5 ~]# pyspark --master spark://spark0:8090
Python 3.8.8 (default, Aug 25 2021, 16:13:02)
[GCC 8.5.0 20210514 (Red Hat 8.5.0-3)] on linux
Type "help", "copyright", "credits" or "license" for more information.
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
23/07/10 09:00:31 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java
classes where applicable
Welcome to

    ___
   / _\ \_  ___ _ _ _/ / _\ \
  _\ \ \_ \_ \_ `/_ / \_ \_ \
 /_ / ._ \_ \_,/_ / / \_ \_ \
  /_/
                                         version 3.3.0

Using Python version 3.8.8 (default, Aug 25 2021 16:13:02)
Spark context Web UI available at http://e5217f917ab5:4040
Spark context available as 'sc' (master = spark://spark0:8090, app id = app-20230710143039-0000).
SparkSession available as 'spark'.
>>> []
```

```
ip="/Software/data/names.json"
op="/Software/data/nameop"
spark.sparkContext.setLogLevel("WARN")
peopleDF = spark.read.json(ip)    # One Job
namesDF = peopleDF.select("firstName","lastName")
namesDF.write.option("header","true").csv(op)  # Second Job
```

```
Spark context available as 'sc' (master = spark://spark0:8090, app id = app-20230710143039-0000).
SparkSession available as 'spark'.
>>> ip="/Software/data/names.json"
>>> op="/Software/data/nameop"
>>> spark.sparkContext.setLogLevel("WARN")
>>> peopleDF = spark.read.json(ip)      # One Job
>>> namesDF = peopleDF.select("firstName", "lastName")
>>> namesDF.write.option("header", "true").csv(op)    # Second Job
>>> █
```

Access the Spark Cluster web console.

<http://127.0.0.1:8080>

You should have two worker nodes as shown below.

Click on Application ID → Application Detail UI

Spark shell - Spark Jobs Application: Spark shell

ID: app-20201113095544-0006  
Name: Spark shell  
User: root  
Cores: Unlimited (2 granted)  
Executor Limit: Unlimited (2 granted)  
Executor Memory: 1024.0 MIB  
Executor Resources:  
Submit Date: 2020/11/13 09:55:44  
State: RUNNING  
[Application Detail UI](#)

▼ Executor Summary (2)

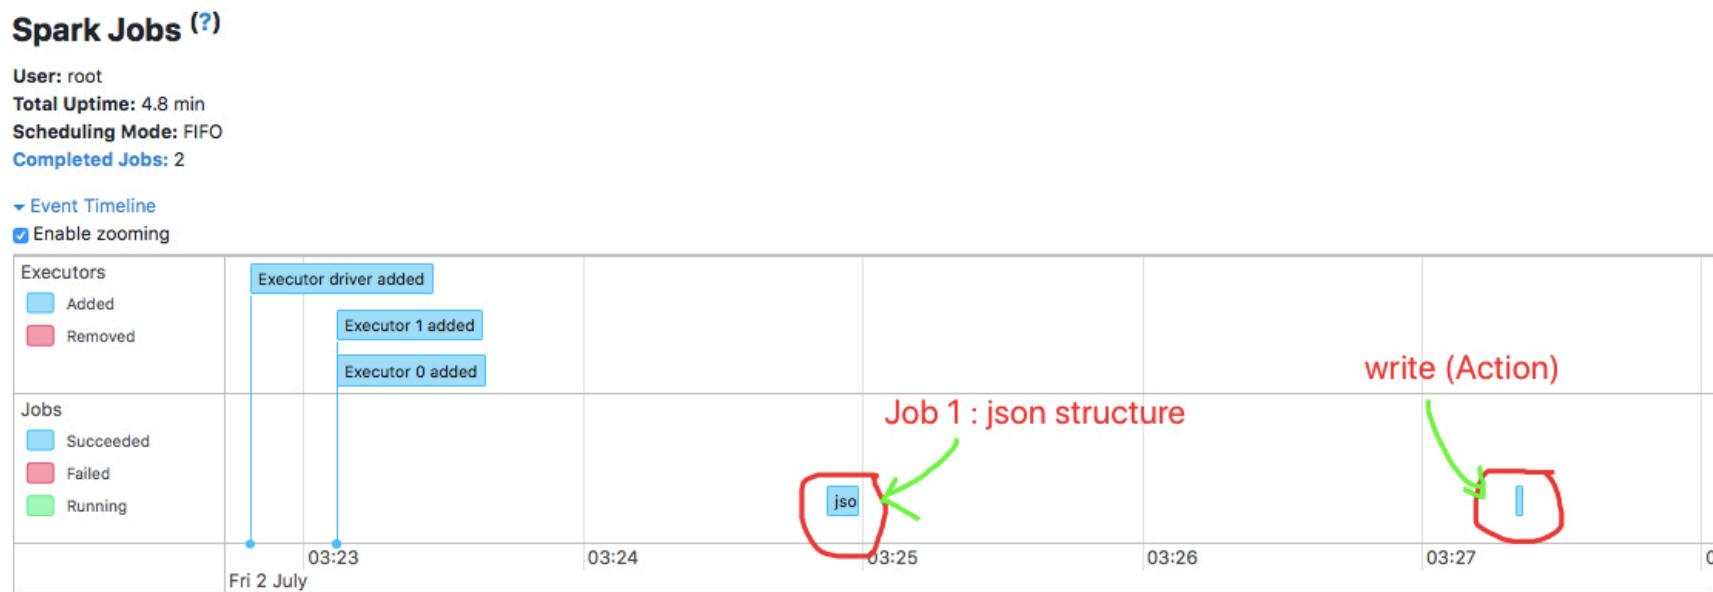
ExecutorID	Worker	Cores	Memory	Resources	State	Logs
1	<a href="#">worker-20201113080544-172.18.0.2-38441</a>	1	1024		RUNNING	<a href="#">stdout stderr</a>
0	<a href="#">worker-20201113080133-172.18.0.3-46519</a>	1	1024		RUNNING	<a href="#">stdout stderr</a>

You can verify the Jobs details Using Spark UI: <http://localhost:4040/jobs/>

Web UI comes with the following tabs (which may not all be visible at once as they are lazily created on demand, e.g. Streaming tab):

- Jobs
- Stages
- Storage with RDD size and memory use
- Environment
- Executors
- SQL/Dataframe

The **Jobs Tab** shows status of all Spark jobs in a Spark application. Expand Eventtime line

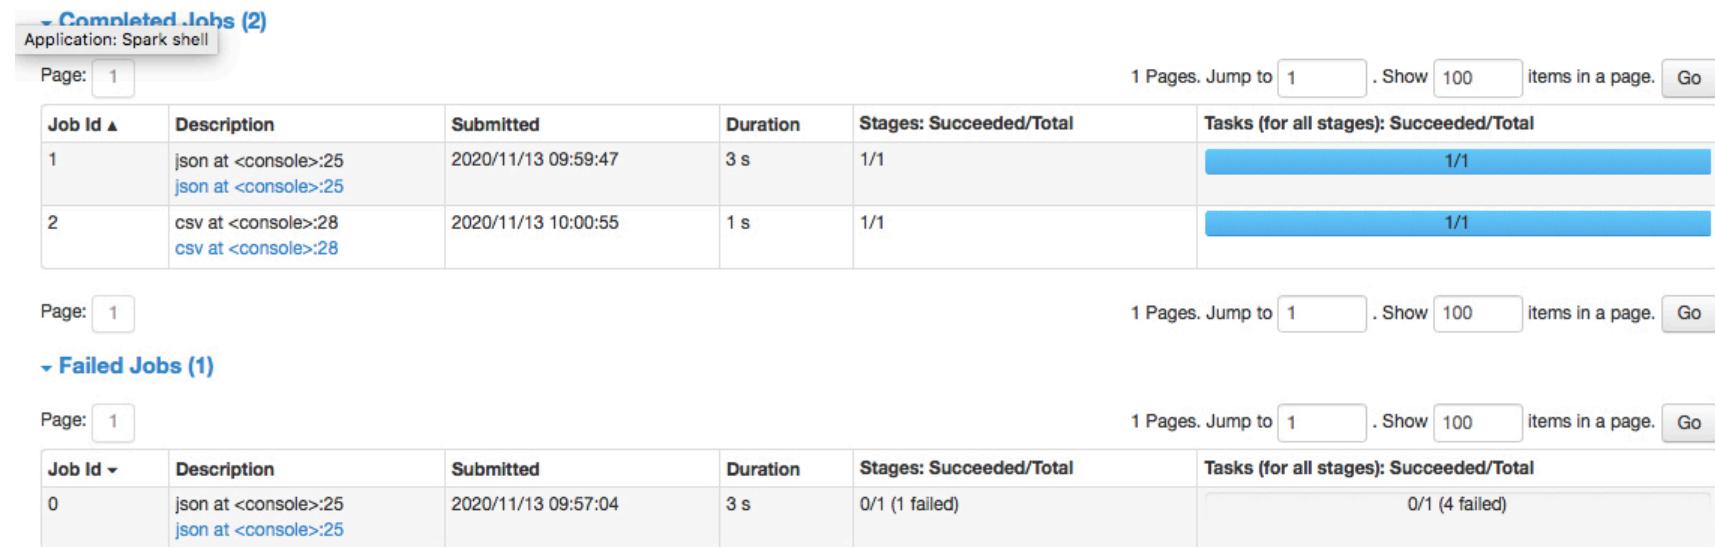


Display the timeline of Executor being added in the Job. When the Job Get completed etc. In the above, you can verify that 2 executors have been added.

- 1 Job Succeeded i.e (Reading Json file)
- 2 Job Succeeded i.e write action.

When you hover over a job in Event Timeline not only you see the job legend but also the job is highlighted in the Summary section.

The Event Timeline section shows not only jobs but also executors. You can verify the completed and failed jobs too from this console:



The screenshot shows two sections of the Apache Spark UI: 'Completed Jobs (2)' and 'Failed Jobs (1)'. Both sections have a table with columns: Job Id, Description, Submitted, Duration, Stages: Succeeded/Total, and Tasks (for all stages): Succeeded/Total.

**Completed Jobs (2) Data:**

Job Id	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
1	json at <console>:25 json at <console>:25	2020/11/13 09:59:47	3 s	1/1	1/1
2	csv at <console>:28 csv at <console>:28	2020/11/13 10:00:55	1 s	1/1	1/1

**Failed Jobs (1) Data:**

Job Id	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
0	json at <console>:25 json at <console>:25	2020/11/13 09:57:04	3 s	0/1 (1 failed)	0/1 (4 failed)

You can verify some of the following Queries from this console?

How long the Job takes to execute? – (Job ID -> 1: Here Duration is 2 s.)

How Many States and Tasks created per Job? ( 2: 1 stages and only one task)

Completed Jobs (2)

Page: 1

1 Pages. Jump to 1 . Show 100 items in a page. Go

Job Id	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
1	csv at NativeMethodAccessorImpl.java:0 csv at NativeMethodAccessorImpl.java:0	2021/07/02 03:27:19	2 s	1/1	1/1
0	json at NativeMethodAccessorImpl.java:0 json at NativeMethodAccessorImpl.java:0	2021/07/02 03:24:52	7 s	1/1	1/1

Click on job (1 – write action or job ) -> for details on Description → stages to view DAG - When you click a job in All Jobs Page page, you see the **Details for Job** page.

The screenshot shows the Apache Spark 3.1.2 UI interface. At the top, there is a navigation bar with tabs: APACHE SPARK 3.1.2, Jobs, Stages (which is highlighted with a red oval), Storage, Environment, Executors, and SQL. Below the navigation bar, the title "Details for Stage 1 (Attempt 0)" is displayed. Under this title, several metrics are listed:

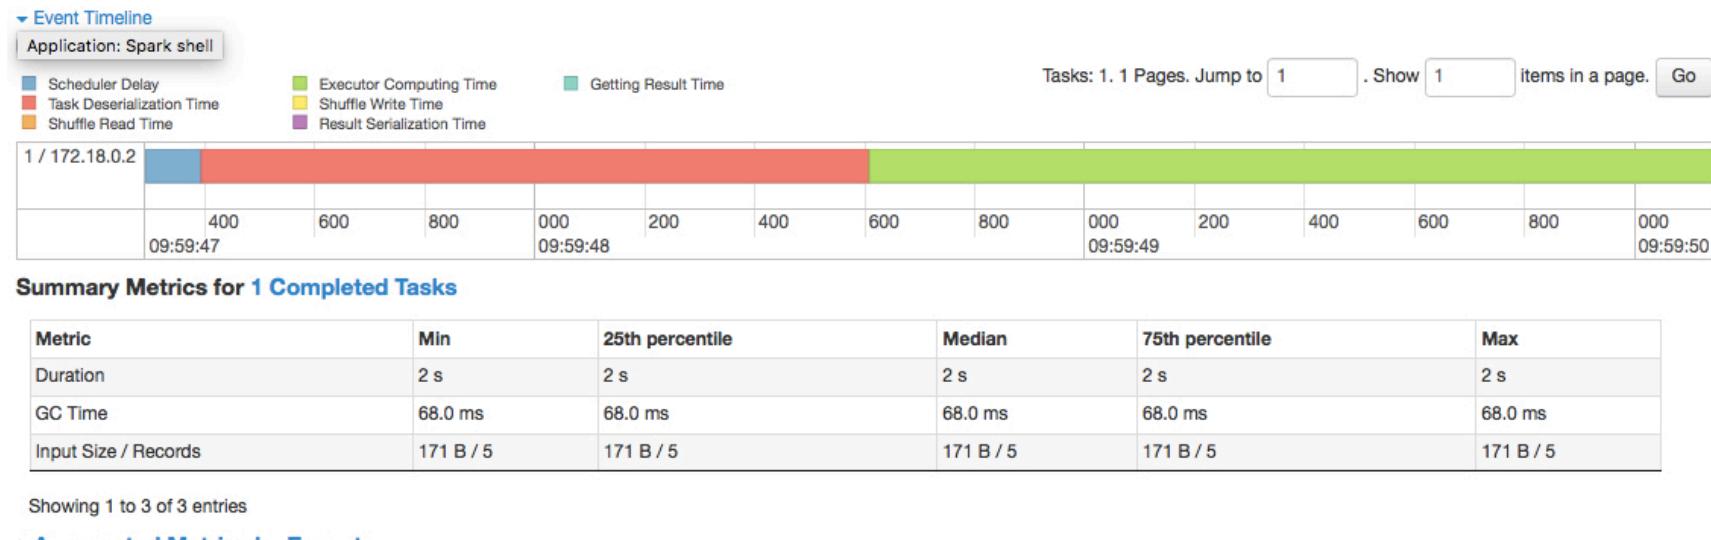
- Resource Profile Id: 0
- Total Time Across All Tasks: 3 s
- Locality Level Summary: Process local: 1
- Input Size / Records: 171.0 B / 4
- Output Size / Records: 73.0 B / 4
- Associated Job Ids: 1

Below these metrics, there is a link labeled "▼ DAG Visualization". A large, rounded rectangular box highlights the DAG visualization area. Inside this box, the title "Stage 1" is at the top. Below it, two blue rounded rectangular boxes represent RDDs. The top box is labeled "Scan json" and contains the text "FileScanRDD [4] csv at NativeMethodAccessorImpl.java:0". An arrow points from this box down to the second box, which is labeled "MapPartitionsRDD [5]" and also contains the text "csv at NativeMethodAccessorImpl.java:0".

As seen above, there is 1 stage in the Job. It scans the file and create a map partition RDD.

It shows the DAG graph and the stage details.

Scroll down to view the Event timeline.



Questions:

How long does the scheduler take to schedule the task? (Check the Light Blue Block)

How long does the task deserialization take? (Check the Orange Block)

What is the actual computation time? (Verify the Green Color block)

In ideal scenario, Green should take most of the computing resources.

You can also verify how much each task took to execute. It will provide some information about the health of the computing task.

Scroll down to view the task metrics: which task take majority of the time etc. How much byte consume or output by each task and its statistics.

**Tasks (1)**

Index	Task ID	Attempt	Status	Locality level	Executor ID	Host	Logs	Launch Time	Duration	GC Time	Input Size / Records	Errors
0	4	0	SUCCESS	PROCESS_LOCAL	1	172.18.0.2	stdout stderr	2020-11-13 15:29:47	2 s	68.0 ms	171 B / 5	

You can verify from the above, Locality\_Level column – the data is process locally. i.e data locality on the same process.

How much the input size for the task above? Hints -> Look for Input Size/ Records column.

## Executors Tab

Stats of each executor, how much time it spent on GC etc. How much data get shuffle?  
 Executors tab in web UI shows

### Summary

	RDD Blocks	Storage Memory	Disk Used	Cores	Active Tasks	Failed Tasks	Complete Tasks	Total Tasks	Task Time (GC Time)	Input	Shuffle Read	Shuffle Write
<b>Active(3)</b>	1	23.4 KB / 1.3 GB	0.0 B	2	0	0	0	0	0 ms (0 ms)	0.0 B	0.0 B	0.0 B
<b>Dead(4)</b>	2	46.8 KB / 1.7 GB	0.0 B	4	0	0	7	7	1.0 min (6 s)	21 KB	0.0 B	0.0 B
<b>Total(7)</b>	3	70.2 KB / 3 GB	0.0 B	6	0	0	7	7	1.0 min (6 s)	21 KB	0.0 B	0.0 B

### Executors

Show 20 entries

Search:

Executor ID	Address	Status	RDD Blocks	Storage Memory	Disk Used	Cores	Active Tasks	Failed Tasks	Complete Tasks	Total Tasks	Task Time (GC Time)			Logs	Thread Dump
											Input	Shuffle Read	Shuffle Write		
driver	192.168.188.173:54112	Active	1	23.4 KB / 434 MB	0.0 B	0	0	0	0	0	0 ms (0 ms)	0.0 B	0.0 B	0.0 B	Thread Dump
0	192.168.188.173:54931	Dead	1	23.4 KB / 434 MB	0.0 B	1	0	0	4	4	37 s (6 s)	9.5 KB	0.0 B	0.0 B	stdout stderr Thread Dump
1	192.168.188.174:34320	Dead	1	23.4 KB / 434 MB	0.0 B	1	0	0	3	3	25 s (0.3 s)	11.5 KB	0.0 B	0.0 B	stdout stderr Thread Dump
2	192.168.188.174:46874	Dead	0	0.0 B / 434 MB	0.0 B	1	0	0	0	0	0 ms (0 ms)	0.0 B	0.0 B	0.0 B	stdout stderr Thread Dump
3	192.168.188.173:34382	Dead	0	0.0 B / 434 MB	0.0 B	1	0	0	0	0	0 ms (0 ms)	0.0 B	0.0 B	0.0 B	stdout stderr Thread Dump
4	192.168.188.174:44548	Active	0	0.0 B / 434 MB	0.0 B	1	0	0	0	0	0 ms (0 ms)	0.0 B	0.0 B	0.0 B	stdout stderr Thread Dump

Click on SQL tab

## Details for Query 0

Submitted Time: 2021/07/02 05:55:43

Duration: 6 s

Succeeded Jobs: 1

Show the Stage ID and Task ID that corresponds to the max metric

### Scan json

number of output rows: 4  
number of files read: 1  
metadata time: 0 ms  
size of files read: 171.0 B

### Execute InsertIntoHadoopFsRelationCommand

number of written files: 1  
written output: 73.0 B  
number of output rows: 4  
number of dynamic part: 0

You can view the statistics of the SQL – How much Rows read and how much bytes etc.

Verify the plan (Scroll down the SQL tab.)

▼ Details

```
== Physical Plan ==
Execute InsertIntoHadoopFsRelationCommand (2)
+- Scan json  (1)

(1) Scan json
Output [2]: [firstName#7, lastName#8]
Batched: false
Location: InMemoryFileIndex [file:/opt/data/names.json]
ReadSchema: struct<firstName:string,lastName:string>

(2) Execute InsertIntoHadoopFsRelationCommand
Input [2]: [firstName#7, lastName#8]
Arguments: file:/opt/data/op, false, CSV, [header=true, path=op], ErrorIfExists, [firstName, lastName]
```

----- Lab Ends Here -----

## 18. Persisting Data – 40 Minutes

In this lab, you will learn how to cache RDD and uncache it, when it's done.

Data being used in this lab:

Create two files in **/Software/data** folder with the following content.

**people.csv**

```
PCODE,lastName,firstName,age
02134,Hopper,Grace,52
94020,Turing,Alan,32
94020,Lovelace,Ada,28
87501,Babbage,Charles,49
02134,Wirth,Niklaus,48
```

**pCodes.csv**

```
PCODE,city,state
02134,Boston,MA
94020,Palo Alto,CA
87501,Santa Fe,NM
60645,Chicago,IL
```

Load the data using dataframe.

Open a terminal and start the **spark-shell** using pyspark. Specify the data path of the file appropriately.

```
over20DF = spark.read.option("header","true").csv("people.csv").where("age > 20")
pcodesDF = spark.read.option("header","true").csv("pcodes.csv")
joinedDF = over20DF.join(pcodesDF, "pcode").persist()
```

Fetch the records which **pincode** is “94020”

```
joinedDF.where("pcode = 94020").show()
```

Fetch another pincode.

```
joinedDF.where("pcode = 87501").show()
```

```

scala> val over20DF = spark.read.option("header","true").csv("people.csv").where("age > 20")
over20DF: org.apache.spark.sql.Dataset[org.apache.spark.sql.Row] = [pcode: string, lastName: string ... 2 more fields]

scala> val pcodesDF = spark.read.option("header","true").csv("pcodes.csv")
pcodesDF: org.apache.spark.sql.DataFrame = [pcode: string, city: string ... 1 more field]

scala> val joinedDF = over20DF.join(pcodesDF, "pcode").persist()
joinedDF: org.apache.spark.sql.Dataset[org.apache.spark.sql.Row] = [pcode: string, lastName: string ... 4 more fields]

scala> 20/11/11 06:19:19 WARN HeartbeatReceiver: Removing executor driver with no recent heartbeats: 137346 ms exceeds timeout 120000 ms
20/11/11 06:19:19 WARN SparkContext: Killing executors is not supported by current scheduler.

scala> joinedDF.where("pcode = 94020").show()
+---+---+---+---+---+
|pcode|lastName|firstName|age|city|state|
+---+---+---+---+---+
|94020| Turing| Alan| 32|Palo Alto| CA|
|94020|Lovelace| Ada| 28|Palo Alto| CA|
+---+---+---+---+---+


scala> joinedDF.where("pcode = 87501").show()
+---+---+---+---+---+
|pcode|lastName|firstName|age|city|state|
+---+---+---+---+---+
|87501| Babbage| Charles| 49|Santa Fe| NM|
+---+---+---+---+---+

```

Register as a table and cache.

```
over20DF.createTempView("overAge")
```

```
spark.sql("CACHE TABLE overAge")
```

```
spark.sql("CACHE TABLE over_age AS SELECT * FROM overAge WHERE age > 20")
spark.sql("select * from over_age").show
```

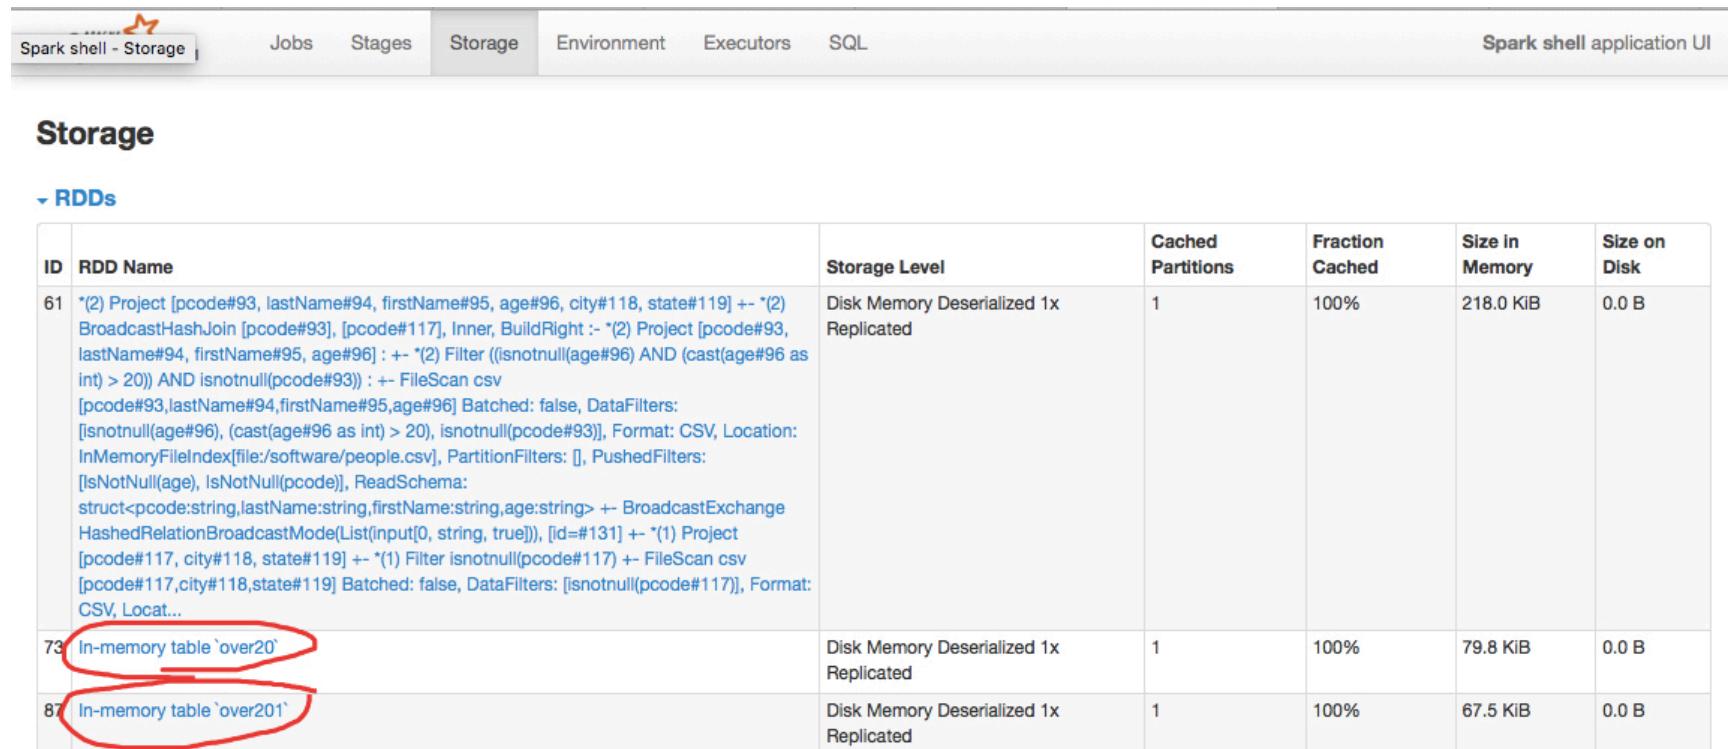
```
scala> over20DF.createTempView("overAge")

scala> spark.sql("CACHE TABLE overAge")
20/11/11 06:50:19 WARN CacheManager: Asked to cache already cached data.
res16: org.apache.spark.sql.DataFrame = []

scala> spark.sql("CACHE TABLE over_age AS SELECT * FROM overAge WHERE age > 20")
20/11/11 06:50:43 WARN CacheManager: Asked to cache already cached data.
res17: org.apache.spark.sql.DataFrame = []

scala> spark.sql("select * from over_age").show
+---+---+---+---+
|lpcode|lastName|firstName|age|
+---+---+---+---+
|102134| Hopper| Grace| 52|
|194020| Turing| Alan| 32|
|194020| Lovelace| Ada| 28|
|187501| Babbage| Charles| 49|
|102134| Wirth| Niklaus| 48|
+---+---+---+---+
```

View the cache details using the URL - <http://localhost:4040/jobs/>



The screenshot shows the Spark shell application UI with the Storage tab selected. The table displays RDD details:

ID	RDD Name	Storage Level	Cached Partitions	Fraction Cached	Size in Memory	Size on Disk
61	<pre>*(2) Project [pcode#93, lastName#94, firstName#95, age#96, city#118, state#119] +- *(2) BroadcastHashJoin [pcode#93], [pcode#117], Inner, BuildRight :- *(2) Project [pcode#93, lastName#94, firstName#95, age#96] : +- *(2) Filter ((isNotNull(age#96) AND (cast(age#96 as int) &gt; 20)) AND isNotNull(pcode#93)) : +- FileScan csv [pcode#93,lastName#94,firstName#95,age#96] Batched: false, DataFilters: [isNotNull(age#96), (cast(age#96 as int) &gt; 20), isNotNull(pcode#93)], Format: CSV, Location: InMemoryFileIndex[file/software/people.csv], PartitionFilters: [], PushedFilters: [IsNotNull(age), IsNotNull(pcode)], ReadSchema: struct&lt;pcode:string,lastName:string,firstName:string,age:string&gt; +- BroadcastExchange HashedRelationBroadcastMode(List(input[0, string, true])), [id=#131] +- *(1) Project [pcode#117, city#118, state#119] +- *(1) Filter isNotNull(pcode#117) +- FileScan csv [pcode#117,city#118,state#119] Batched: false, DataFilters: [isNotNull(pcode#117)], Format: CSV, Locat...</pre>	Disk Memory Deserialized 1x Replicated	1	100%	218.0 KIB	0.0 B
73	In-memory table `over20`	Disk Memory Deserialized 1x Replicated	1	100%	79.8 KIB	0.0 B
87	In-memory table `over20`	Disk Memory Deserialized 1x Replicated	1	100%	67.5 KIB	0.0 B

Changing the **storage** tab. Verify the size in memory.

Let us change the storage type.

Specify the Storage location as Disk.

```
import org.apache.spark.storage.StorageLevel  
joinedDF.persist(StorageLevel.DISK_ONLY)
```

Remove the cache and cache again to change the persist level.

```
joinedDF.unpersist()  
joinedDF.cache()
```

```
scala>

scala> import org.apache.spark.storage.StorageLevel
import org.apache.spark.storage.StorageLevel

scala> joinedDF.persist(StorageLevel.DISK_ONLY)
20/11/11 06:56:26 WARN CacheManager: Asked to cache already cached data.
res19: joinedDF.type = [PCODE: string, LastName: string ... 4 more fields]

scala> joinedDF.unpersist()
res20: joinedDF.type = [PCODE: string, LastName: string ... 4 more fields]

scala> joinedDF.persist(StorageLevel.DISK_ONLY)
res21: joinedDF.type = [PCODE: string, LastName: string ... 4 more fields]

scala> joinedDF.cache()
20/11/11 06:57:39 WARN CacheManager: Asked to cache already cached data.
res22: joinedDF.type = [PCODE: string, LastName: string ... 4 more fields]

scala>
```

To Remove the table/view – over20 from cache follow the steps given below.  
From the web UI – storage section. You can verify the cache details. Currently there are two objects in cache.

Before removing the cache:

**Storage**

**RDDs**

ID	RDD Name	Storage Level	Cached Partitions	Fraction Cached	Size in Memory	Size on Disk
73	In-memory table `over20`	Disk Memory Deserialized 1x Replicated	1	100%	99.3 KIB	0.0 B
87	In-memory table `over201`	Disk Memory Deserialized 1x Replicated	1	100%	87.0 KIB	0.0 B

Display the list of tables:

`spark.catalog.listTables.show`

`spark.catalog.uncacheTable("overage")`

After executing the above command: the above mention cache should be removed as shown below:

**Storage**

**RDDs**

ID	RDD Name	Storage Level	Cached Partitions	Fraction Cached	Size in Memory	Size on Disk
87	In-memory table `over201`	Disk Memory Deserialized 1x Replicated	1	100%	160.7 KiB	0.0 B

-----Lab Ends Here-----

## 19. Broadcast & Repartition – 45 Minutes

Open a spark shell terminal and perform the following steps.

In this we are broadcasting the states and countries details to all the executor.

### Understanding Broadcast.

#pyspark

Perform the following steps on the spark shell.

Import the necessary packages.

```
from pyspark.sql.functions import broadcast
```

Define the States and Countries data set.

```
states = [("NY","New York"),("CA","California"),("FL","Florida"),("MH","Mahrashtra")]
countries = [("USA","United States of America"),("IN","India")]
```

Convert the state and countries information into Dataframe.

```
stDF = spark.createDataFrame(data=states, schema = ["cs","state"])
csDF = spark.createDataFrame(data=countries, schema = ["cc","country"])
```

Define a dataset of person details.

```
data = [("James","Smith","USA","CA"), \
        ("Michael","Rose","USA","NY"), \
        ("Robert","Williams","USA","CA"), \
        ("Maria","Jones","USA","FL"),("Ram","Kumar","IN","MH")]
```

Define the column structure of the data frame.

```
columns = ["firstname","lastname","country","state"]
```

Person Dataframe construction with schema.

```
perDF = spark.createDataFrame(data=data, schema = columns)
perDF.printSchema()
perDF.show(truncate=False)
```

```
In [59]: perDF = spark.createDataFrame(data=data, schema = columns)
perDF.printSchema()
perDF.show(truncate=False)
```

```
root
 |-- firstname: string (nullable = true)
 |-- lastname: string (nullable = true)
 |-- cc: string (nullable = true)
 |-- sc: string (nullable = true)

+-----+-----+---+
|firstname|lastname|cc |sc |
+-----+-----+---+
|James    |Smith    |USA|CA |
|Michael  |Rose     |USA|NY |
|Robert   |Williams|USA|CA |
|Maria    |Jones    |USA|FL |
|Ram      |Kumar    |IN |MH |
+-----+-----+---+
```

Get the Country name for each state using the country code Using broadcast.

```
result = perDF.join(broadcast(csDF), perDF["cc"] == csDF["cc"])
result.show()
```

Derive the country from the code using broadcast

```
In [60]: result = perDF.join(broadcast(csDF), perDF["cc"] == csDF["cc"])
```

```
In [67]: result.show()
```

firstname	lastname	cc	sc	cc	country	cs	state
James	Smith	USA	CA	USA	United States of ...	CA	California
Michael	Rose	USA	NY	USA	United States of ...	NY	New York
Robert	Williams	USA	CA	USA	United States of ...	CA	California
Maria	Jones	USA	FL	USA	United States of ...	FL	Florida
Ram	Kumar	IN	MH	IN	India	MH	Mahrashtra

Derive the State from the code using broadcast.

```
result = result.join(broadcast(stDF), result["sc"] == stDF["cs"])
result.show()
```

```
In [62]: result = result.join(broadcast(stDF), result["sc"] == stDF["cs"])
```

```
In [63]: result.show()
```

firstname	lastname	cc	sc	cc	country	cs	state
James	Smith	USA	CA	USA	United States of ...	CA	California
Michael	Rose	USA	NY	USA	United States of ...	NY	New York
Robert	Williams	USA	CA	USA	United States of ...	CA	California
Maria	Jones	USA	FL	USA	United States of ...	FL	Florida
Ram	Kumar	IN	MH	IN	India	MH	Mahrastrha

Drop the code - state and country from the Dataframe.

```
result.drop("cc","sc","cs").show()
```

```
In [64]: result.drop("cc","sc","cs").show()
```

firstname	lastname	country	state
James	Smith	United States of ...	California
Michael	Rose	United States of ...	New York
Robert	Williams	United States of ...	California
Maria	Jones	United States of ...	Florida
Ram	Kumar	India	Mahrasthra

## Understanding Repartition

Create a simple RDD using the `sc.parallelize` method to determine the number of partitions spark creates by default.

```
data = [1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11]
```

```
rdd = sc.parallelize(data)
print(rdd.getNumPartitions())
```

```
In [21]: data = [1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11]
```

```
In [22]: rdd = sc.parallelize(data)
```

```
In [23]: print(rdd.getNumPartitions())
```

```
6
```

As shown above, spark creates 6 partitions by default here (Number of cores of your machine). Now, you must be wondering where this default number has come from.

Spark has a default parallelism parameter which is determined by, `sc.defaultParallelism`. When we run this method, it returns 6 as shown below, which is having 6 cores. It may be different in your case, depending on your machine cores.

### sc.defaultParallelism

This is how the data is present inside each partition. Spark uses an internal Hash Partitioning Scheme to split the data into these smaller chunks.

We can use the `rdd.glom` method to display the partitions in a list.

`glom` - returns an RDD having the elements within each partition in a separate list

```
# rdd.glom().collect()
```

In [24]: `sc.defaultParallelism`

[Reload this page](#)

Out[24]: 6

In [25]: `rdd.glom().collect()`

Out[25]: [[1], [2, 3], [4, 5], [6, 7], [8, 9], [10, 11]]

In above, Inner Array represents the **Partition**. As shown above it represent 6 partitions.

Let's visualize the above collect operation in the Spark WebUI(if you are using spark locally, navigate to <http://localhost:4040> to see the spark webui or if spark is being accessed via a cluster navigate to your cluster specific localhost webUI)

Note: Spark shell notifies a port number before creating a Sparksession.

Since 2 partitions are present, 2 tasks would be launched for this action (Depends if you have that cores or executor), as shown below,

▼ Completed jobs (3)

[Reload this page](#) 1 Pages. Jump to  . Show  items in a page. [Go](#)

Job Id ▾	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
2	collect at /tmp/ipykernel_113/27051842.py:1 collect at /tmp/ipykernel_113/27051842.py:1	2023/08/21 16:35:55	57 ms	1/1	6/6
1	collect at /tmp/ipykernel_113/27051842.py:1 collect at /tmp/ipykernel_113/27051842.py:1	2023/08/21 16:35:25	85 ms	1/1	6/6
0	collect at /tmp/ipykernel_113/27051842.py:1 collect at /tmp/ipykernel_113/27051842.py:1	2023/08/21 16:34:57	0.4 s	1/1	2/2

How to see number of Partitions in a Dataframe?

The above method used for an RDD can also be applied to a dataframe. Let's convert the RDD to a Dataframe and apply the same method. We can see that the same number of partitions are present.

```
df = rdd.toDF("Integer")
```

Converting rdd to Dataframe

```
df.rdd.glom().collect()
```

```
In [28]: df = rdd.toDF("Integer")
```

```
In [30]: df.rdd.glom().collect()
```

```
Out[30]: [[Row(value=1)],
           [Row(value=2), Row(value=3)],
           [Row(value=4), Row(value=5)],
           [Row(value=6), Row(value=7)],
           [Row(value=8), Row(value=9)],
           [Row(value=10), Row(value=11)]]
```

By Default, partition is divided by size in case of file.

Now, you must've had the question in your mind as to how spark partitions the data when reading textfiles.

Let's read a simple csvfile and see the number of partitions here.

```
peopleDF = spark.read.format("csv").option("sep", ",").option("inferSchema", "true").option("header", "true").load("/Software/data/zcodes.csv")
```

```
peopleDF.rdd.getNumPartitions
```

```
In [32]: peopleDF = spark.read.format("csv").option("sep", ",") \
.option("inferSchema", "true").option("header", "true").load("/Software/data/zcodes.csv")
```

```
In [34]: peopleDF.rdd.getNumPartitions()
```

```
Out[34]: 1
```

We can see that only 1 partition is created here. Alright let's break this down,  
*Spark by default creates 1 partition for every 128 MB of the file.*

So if you are reading a file of size 1GB, it creates 10 partitions.

**So how to control the number of bytes a single partition can hold ?**

The file being read here is less than that of the minimum size, hence only 1 partition is created in this case.

The no. of partitions is determined by `spark.sql.files.maxPartitionBytes` parameter, which is set to 128 MB, by default.

```
spark.conf.get("spark.sql.files.maxPartitionBytes") # -> 128 MB by default
```

```
-----
```

```
In [35]: spark.conf.get("spark.sql.files.maxPartitionBytes")
```

```
Out[35]: '134217728b'
```

**Note:** The files being read must be splitable by default for spark to create partitions when reading the file. So, in case of compressed files like snappy, gz or lzo etc, a single partition is created irrespective of the size of the file.

Let us upload a file of larger size.

```
sfDF = spark.read.format("csv").option("sep", ",").option("inferSchema",  
"true").option("header", "true").load("/Software/data/departuredelays.csv")
```

```
sfDF.rdd.getNumPartitions()
```

```
In [37]: sfDF = spark.read.format("csv").option("sep", ",") \
.option("inferSchema", "true").option("header", "true").load("/Software/data/departuredelays.csv")
```

```
In [38]: sfDF.rdd.getNumPartitions()
```

Reload this page

```
Out[38]: 6
```

Let's manually set the `spark.sql.files.maxPartitionBytes` to 4MB, so that we get 9 partitions upon reading the file.

```
spark.conf.set("spark.sql.files.maxPartitionBytes", 4000000)
spark.conf.get("spark.sql.files.maxPartitionBytes")
```

Upload the file again

```
sfDF1 = spark.read.format("csv").option("sep", ",").option("inferSchema",
"true").option("header", "true").load("/Software/data/departuredelays.csv")
sfDF1.rdd.getNumPartitions()
```

Now, there are 9 Partitions. It depends on the size of your file.

```
In [39]: spark.conf.set("spark.sql.files.maxPartitionBytes", 4000000)
spark.conf.get("spark.sql.files.maxPartitionBytes")
Out[39]: '4000000' Reload this page

In [40]: sfDF1 = spark.read.format("csv").option("sep", ",").option("inferSchema", "true").option("header", "true").load("/S
sfDF1.rdd.getNumPartitions()
Out[40]: 9
```

Let's analyse with a simple example where we create an RDD with a single partition and perform an action on the same.

When the above action is seen on the Spark WebUI, only 2 tasks would be assigned to process this data. So, imagine a case where you are processing a huge volume of data without the concept of partitioning, then the entire data would be processed by a single executor taking a lot of time and memory.

The screenshot shows the Apache Spark 3.0.1 Web UI interface. At the top, there is a navigation bar with links for Jobs, Stages, Storage, Environment, Executors, and SQL. On the right side of the header, it says "Spark shell application UI". Below the header, the title "Spark Jobs" is displayed with a help icon. Underneath, there is some system information: User: root, Total Uptime: 1.1 h, Scheduling Mode: FIFO, and Completed Jobs: 15. There are two links: "Event Timeline" and "Completed Jobs (15)". The main content area shows a table of completed jobs. The table has columns: Job Id, Description, Submitted, Duration, Stages: Succeeded/Total, and Tasks (for all stages): Succeeded/Total. One specific row is highlighted with a red oval around the "Tasks (for all stages): Succeeded/Total" cell, which contains "2/2". The table also includes pagination controls at the bottom: "Page: 1", "1 Pages. Jump to", "Show 100 items in a page.", and a "Go" button.

Job Id	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
14	collect at <console>:29 collect at <console>:29	2020/11/21 07:36:42	27 ms	1/1	2/2

Now what if you wish to control these partitions by yourself. This is where the methods of `repartition` and `coalesce` come to effect.

### Repartition and Coalesce

Create a data collection.

```
mydata = [1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12]
```

Convert the collection into RDD with one Partition only.

```
rdd1 = sc.parallelize(mydata,1)
```

Determine the number of partitions. It should be only 1.

```
rdd1.getNumPartitions()
```

## Repartition and Coalesce

```
In [2]: mydata = [1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12]
```

```
In [6]: rdd1 = sc.parallelize(mydata, 1)
```

```
In [9]: rdd1.getNumPartitions()
```

```
Out[9]: 1
```

Let us convert RDD to data frame for further analysis.

```
df = rdd1.toDF("Integer")
```

Repartition the above dataframe to 10 partitions from 1.

```
partDF = df.repartition(10)
```

```
partDF.rdd.getNumPartitions()
```

Now the number of partitions should be 10.

```
In [11]: df = rdd1.toDF("Integer")  
In [12]: partDF = df.repartition(10)  
In [13]: partDF.rdd.getNumPartitions()  
Out[13]: 10
```

---

After repartitioning, the partition size has increased from 1 to 10 as shown above.

You can observe the partitions information or data as shown below.

```
spark.conf.get("spark.sql.shuffle.partitions")  
partDF.rdd.glom().collect()
```

```
In [14]: spark.conf.get("spark.sql.shuffle.partitions")
```

```
Out[14]: '200'
```

```
In [16]: partDF.rdd.glom().collect()
```

```
Out[16]: [[Row(value=4)],
           [Row(value=11), Row(value=1)],
           [Row(value=3), Row(value=12)],
           [Row(value=7)],
           [Row(value=10)],
           [Row(value=5)],
           [Row(value=6)],
           [Row(value=8)],
           [Row(value=2)],
           [Row(value=9)]]
```

Let's visualise the same on the Spark WebUI

```
dfre = df.repartition(3)
dfre.rdd.glom().collect()
```

```
In [17]: dfre = df.repartition(3)
```

```
In [18]: dfre.rdd.glom().collect()
```

```
Out[18]: [[Row(value=7), Row(value=6), Row(value=9), Row(value=12)],
           [Row(value=11), Row(value=10), Row(value=8), Row(value=4)],
           [Row(value=3), Row(value=5), Row(value=2), Row(value=1)]]
```

The above execution when seen on the spark web console would show 3 tasks being issued on the collect operation(stage 2).

Completed Jobs (18)						
Job Id	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total	
17	collect at <console>:26 collect at <console>:26	2023/05/20 09:58:32	60 ms	1/1 (1 skipped)	3/3 (1 skipped)	
16	rdd at <console>:26 rrd at <console>:26	2023/05/20 09:58:32	12 ms	1/1	1/1	
15	collect at <console>:26 collect at <console>:26	2023/05/20 09:57:42	0.1 s	1/1 (1 skipped)	10/10 (1 skipped)	

Click on the Job description of Job ID 17 – It may be different for your case.

**Details for Job 17****Status:** SUCCEEDED**Submitted:** 2023/05/20 09:58:32**Duration:** 60 ms**Completed Stages:** 1**Skipped Stages:** 1

- ▶ Event Timeline
- ▶ DAG Visualization

▼ **Completed Stages (1)**Page: 1 Pages. Jump to  . Show  items in a page. 

Stage Id ▾	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
19	collect at <console>:26 +details	2023/05/20 09:58:32	42 ms	3/3			258.0 B	

Page: 1 Pages. Jump to  . Show  items in a page. 

As shown above, repartition causes shuffling of data.

**Coalesce**

Unlike repartition, **coalesce** doesn't perform a shuffle to create the partitions.

Let's analyse this using our dataset, let's reduce the number of partitions to 2 now,

```
partDF.rdd.glom().collect()
partDF.rdd.getNumPartitions()
```

Currently it has 10 partitions, let us reduce to 2.

```
part3 = partDF.coalesce(2)
part3.rdd.getNumPartitions()
```

```
In [21]: part3 = partDF.coalesce(2)
part3.rdd.getNumPartitions()
```

```
Out[21]: 2
```

[https://datanoon.com/blog/spark\\_repartition\\_coalesce/](https://datanoon.com/blog/spark_repartition_coalesce/)

<https://medium.com/@mrpowers/managing-spark-partitions-with-coalesce-and-repartition-4050c57ad5c4>

## 20. Spark – Hive Integration – 45 Minutes

When working with Hive, one must instantiate SparkSession with Hive support, including connectivity to a persistent Hive metastore, support for Hive serdes, and Hive user-defined functions. Users who do not have an existing Hive deployment can still enable Hive support.

use spark.sql.warehouse.dir to specify the default location of database in warehouse

Open a pyspark terminal (Using CLI only)

```
#pyspark --conf spark.sql.catalogImplementation=hive
```

Execute the following code.

```
from os.path import abspath

from pyspark.sql import SparkSession
from pyspark.sql import Row

# warehouse_location points to the default location for managed databases
# and tables
warehouse_location = abspath('/opt/spark/spark-warehouse')
```

Only for CLI Mode.

```
spark = SparkSession \
    .builder \
    .appName("Python Spark SQL Hive integration example") \
    .config("spark.sql.warehouse.dir", warehouse_location) \
    .config("spark.sql.catalogImplementation", "hive") \
    .enableHiveSupport() \
    .getOrCreate()
```

Create a hive datawarehouse table.

```
spark.sql("CREATE TABLE IF NOT EXISTS src (key INT, value STRING) USING
hive")
```

Load the data in the hive table, **src**

```
spark.sql("LOAD DATA LOCAL INPATH
'/opt/spark/examples/src/main/resources/kv1.txt' INTO TABLE src")
```

Some Queries to retrieve data from the `src` table.

```
# Queries are expressed in HiveQL  
spark.sql("SELECT * FROM src").show()
```

```
# +---+-----+  
# |key| value|  
# +---+-----+  
# |238|val_238|  
# | 86| val_86|  
# |311|val_311|  
# ...
```

```
# Aggregation queries are also supported.  
spark.sql("SELECT COUNT(*) FROM src").show()
```

```
# +-----+  
# |count(1)|  
# +-----+  
# |      500 |  
# +-----+
```

```
# The results of SQL queries are themselves DataFrames and support all  
normal functions.  
sqlDF = spark.sql("SELECT key, value FROM src WHERE key < 10 ORDER BY  
key")
```

```
# The items in DataFrames are of type Row, which allows you to access each  
column by ordinal.
```

```
stringsDS = sqlDF.rdd.map(lambda row: "Key: %d, Value: %s" % (row.key,  
row.value))  
for record in stringsDS.collect():  
    print(record)
```

```
# Key: 0, Value: val_0  
# Key: 0, Value: val_0  
# Key: 0, Value: val_0  
# ...
```

```
# You can also use DataFrames to create temporary views within a  
SparkSession.
```

```
Record = Row("key", "value")
recordsDF = spark.createDataFrame([Record(i, "val_" + str(i)) for i in range(1,
101)])
recordsDF.createOrReplaceTempView("records")
# Queries can then join DataFrame data with data stored in Hive.
spark.sql("SELECT * FROM records r JOIN src s ON r.key = s.key").show()
```

Execution Output:



```
>>> spark.sql("LOAD DATA LOCAL INPATH '/opt/spark/examples/src/main/resources/kv1.txt' INTO TABLE src")
21/07/03 10:11:29 WARN ObjectStore: Failed to get database global_temp, returning NoSuchObjectException
DataFrame[]

>>> spark.sql("SELECT * FROM src").show()
+---+---+
|key| value|
+---+---+
|238|val_238|
| 86| val_86|
|311|val_311|
| 27| val_27|
|165|val_165|
|409|val_409|
|255|val_255|
|278|val_278|
| 98| val_98|
```

```
>>> spark.sql("SELECT COUNT(*) FROM src").show()
+-----+
| count(1)|
+-----+
|      500|
+-----+

>>> sqlDF = spark.sql("SELECT key, value FROM src WHERE key < 10 ORDER BY key")
>>> stringsDS = sqlDF.rdd.map(lambda row: "Key: %d, Value: %s" % (row.key, row.value))
>>> for record in stringsDS.collect():
...     print(record)
...
Key: 0, Value: val_0
Key: 0, Value: val_0
Key: 0, Value: val_0
Key: 2, Value: val_2
Key: 4, Value: val_4
Key: 5, Value: val_5
Key: 5, Value: val_5
Key: 5, Value: val_5
```

```
key, value, val_2
>>> Record = Row("key", "value")
>>> recordsDF = spark.createDataFrame([Record(i, "val_" + str(i)) for i in range(1, 101)])
>>> recordsDF.createOrReplaceTempView("records")
>>> spark.sql("SELECT * FROM records r JOIN src s ON r.key = s.key").show()
+---+---+---+
|key|value|key|value|
+---+---+---+
| 2|val_2| 2|val_2|
| 4|val_4| 4|val_4|
| 5|val_5| 5|val_5|
| 5|val_5| 5|val_5|
| 5|val_5| 5|val_5|
| 8|val_8| 8|val_8|
| 9|val_9| 9|val_9|
|10|val_10|10|val_10|
|11|val_11|11|val_11|
|12|val_12|12|val_12|
|12|val_12|12|val_12|
|15|val_15|15|val_15|
|15|val_15|15|val_15|
```

----- Lab Ends Here -----

## 21. Spark integration with Hive Cluster (Local Mode)

Pre requisites: Spark Installation and Hive Local Mode installation.

You can refer Hive Lab for Hive Local Installation.

In this lab, we will configure to use Apache Spark with Hive.

```
export SPARK_HOME=/opt/spark  
export HIVE_HOME=/opt/hive_local
```

Set HIVE\_HOME and SPARK\_HOME accordingly.

## 22. Annexure & Errata:

Quit Scala CLI :- :q

Error: No suitable device found: no device found for connection 'System etho'.

You could try removing/renaming /etc/udev/rules.d/70-persistent-net.rules and reboot. A new file will be created, sometimes that fixes this kind of problems.

You could also add a file "/etc/sysconfig/network-scripts/ifcfg-eth1", can you bring up eth1 then?

Changing hostname:

- /etc/hosts
- vi /etc/sysconfig/network
- sysctl kernel.hostname=slave
- bash

Unable to start spark shell with the following error

```
at org.apache.hadoop.hive.ql.session.SessionState.createSessionDirs(SessionState.java:554)
at org.apache.hadoop.hive.ql.session.SessionState.start(SessionState.java:508)
... 85 more
```

Caused by: java.net.ConnectException: Connection refused

```
at sun.nio.ch.SocketChannelImpl.checkConnect(Native Method)
at sun.nio.ch.SocketChannelImpl.finishConnect(SocketChannelImpl.java:717)
at org.apache.hadoop.net.SocketIOWithTimeout.connect(SocketIOWithTimeout.java:206)
at org.apache.hadoop.net.NetUtils.connect(NetUtils.java:531)
at org.apache.hadoop.net.NetUtils.connect(NetUtils.java:495)
at org.apache.hadoop.ipc.Client$Connection.setupConnection(Client.java:614)
at org.apache.hadoop.ipc.Client$Connection.setupIOstreams(Client.java:712)
```

```
at org.apache.hadoop.ipc.Client$Connection.access$2900(Client.java:375)
at org.apache.hadoop.ipc.Client.getConnection(Client.java:1528)
at org.apache.hadoop.ipc.Client.call(Client.java:1451)
```

Solution: unset HADOOP\_CONF\_DIR and unset HADOOP\_HOME

<https://jaceklaskowski.gitbooks.io/mastering-apache-spark/>

issue: Cannot assign requested address

```
17/03/15 21:44:15 ERROR spark.SparkContext: Error initializing SparkContext.
java.net.BindException: Cannot assign requested address: Service 'sparkDriver' failed after 16 retries (starting from 0)! Consider explicitly setting the appropriate port for the service 'sparkDriver' (for example spark.ui.port for SparkUI) to an available port or increasing spark.port.maxRetries.
```

```
at sun.nio.ch.Net.bind0(Native Method)
at sun.nio.ch.Net.bind(Net.java:433)
at sun.nio.ch.Net.bind(Net.java:425)
at sun.nio.ch.ServerSocketChannelImpl.bind(ServerSocketChannelImpl.java:223)
```

Solution:

- include host to ip entry in /etc/hosts --> 192.168.188.178 master
- update in conf/spark-env.sh
  - SPARK\_LOCAL\_IP=127.0.0.1

- SPARK\_MASTER\_HOST=192.168.188.178
- comment the following in conf/spark-defaults.conf
  - # spark.master spark://master:7077

Unable to start spark shell:

Caused by: java.io.IOException: Error accessing /opt/jdk/jre/lib/ext/.\_cld\*.jar

Such kind of error resolve by removing all the hidden file. .\*(You can determine the file by running #ls -alt). List all the hidden dot file and remove it.

Yum repo config

```
# mkdir -p /redhatimg
# mount -o loop,ro /mnt/hgfs/MyExperiment/rhel-server-6.4-x86_64-dvd.iso /redhatimg
```

Create an entry in /etc/fstab so that the system always mounts the DVD image after a reboot.  
 /mnt/hgfs/MyExperiment/rhel-server-6.4-x86\_64-dvd.iso /redhatimg iso9660 loop,ro o o

/etc/yum.repos.d/rheldiso.repo

```
[rhel6dvdiso]
name=RedHatOS DVD ISO
mediaid=1359576196.686790
baseurl=file:///redhatimg
enabled=1
gpgcheck=1
gpgkey=file:///etc/pki/rpm-gpg/RPM-GPG-KEY-redhat-release
```

## ***Issue in Pyspark: TypeError: Too many parameters for typing.Iterable; actual 2, expected 1***

### **Background:**

Python 3.6.8

3.3.0

File "/usr/local/lib/python3.6/site-packages/typing\_extensions.py", line 113, in \_check\_generic  
    raise TypeError(f"Too {'many' if alen > elen else 'few'} parameters for {cls};"  
TypeError: Too many parameters for typing.Iterable; actual 2, expected 1

### **Solution:**

Check python version :

```
#python3 -V
```

```
#yum install python3.8
# python3.8 --version
```

Let's set default python3 to python3.8 on RHEL 8:

Perform the following steps on your machine

```
sudo alternatives --config python3
```

or

```
sudo update-alternatives --config python3
```

```
[root@spark0 /]# alternatives --config python3

There are 2 programs which provide 'python3'.

Selection    Command
-----
*+ 1          /usr/bin/python3.6
              /usr/bin/python3.8

Enter to keep the current selection[+], or type selection number: 2
[root@spark0 /]#
```

Accept 2 as atternatives.

Issue: Traceback (most recent call last):

File "/usr/local/lib/python3.6/site-packages/notebook/notebookapp.py", line 2027, in  
init\_server\_extensions

`python3.8 -m pip install jupyter_contrib_nbextensions`

## Bash profile

```
export JAVA_HOME=/opt/jdk
export PATH=$PATH:$JAVA_HOME/bin
export PATH=$PATH:/opt/spark/bin:/Software/maven/bin
export PYSPARK_PYTHON=python3.8
export SPARK_CONF_DIR=/opt/spark
export SPARK_HOME=/opt/spark
export PYSPARK_DRIVER_PYTHON=jupyter
#export PYSPARK_DRIVER_PYTHON=python3.8
#export PYSPARK_SUBMIT_ARGS='--packages org.apache.spark:spark-sql-kafka-o-
10_2.12:3.3.0,org.apache.spark:spark-streaming-ka
fka-o-10_2.12:3.3.0 pyspark'
export PYSPARK_DRIVER_PYTHON_OPTS='notebook'

export PACKAGES="io.delta:delta-core_2.12:2.1.0",org.apache.spark:spark-sql-kafka-o-10_2.12:3.3.0
#export PYSPARK_SUBMIT_ARGS="--packages ${PACKAGES} pyspark-shell"
export PYSPARK_PYTHON=python3.8
# This is for spylon jupyter integration required for pyspark in Delta.
alias snotebook='pyspark --packages io.delta:delta-core_2.12:2.1.0 --conf
spark.sql.extensions=io.delta.sql.DeltaSparkSessio
nExtension --conf "spark.sql.catalog.spark_catalog=org.apache.spark.sql.delta.catalog.DeltaCatalog" --
conf spark.sql.warehouse
se.dir=/delta/spark-warehouse'

#Required for pyspark+kafka integration.
alias pnotebook='pyspark --packages org.apache.spark:spark-sql-kafka-o-
10_2.12:3.3.0,org.apache.spark:spark-streaming-kafka-
```

```
0-10_2.12:3.3.0 --conf spark.sql.catalogImplementation=hive'
```